

# Optimizers for Deep Neural Networks

Yang You

Presidential Young Professor at NUS Computer Science

[ai.comp.nus.edu.sg](http://ai.comp.nus.edu.sg)

- Typical Assumption: Loss Function is L-Smooth

$$\|\nabla L(\theta_2; \xi) - \nabla L(\theta_1; \xi)\|_2 \leq L \|\theta_2 - \theta_1\|_2 \quad \forall \theta_1, \theta_2, \xi$$

- Empirical risk minimization: draw samples  $\xi$  from a distribution  $\mathbb{P}$

$$\nabla L(\theta_t) = \nabla \mathbb{E}[L(\theta_t; \xi)]$$

- Updating rule for SGD

$$\theta_t = \theta_{t-1} - \eta_t * \nabla L(\theta_{t-1})$$

- $\eta_t$  is a number (learning rate)
- $\theta_t$  is a vector (weights or parameters)
- only uses a batch of data to compute gradients each iteration

# AdaGrad (Adaptive Gradient) Duchi et al. 2011

$$g_t = \nabla L(\theta_{t-1})$$

$$v_t = v_{t-1} + g_t \odot g_t$$

$$\theta_t = \theta_{t-1} - \eta_t * \frac{g_t}{\sqrt{v_t + \epsilon}}$$

- $\theta_t, g_t, v_t$  is a vector;  $\eta_t$  is a number (default value is 0.01)
- It eliminates the need to manually tune the learning rate
- Sparse gradient has a larger update (element-wise updating) and moves faster
- Accumulation of squared gradients in the denominator (**weakness**):
  - Since every added term is positive, the accumulated sum keeps growing during training
  - This in turn causes the update to shrink and eventually become infinitesimally small, at which point the algorithm is no longer able to acquire additional knowledge

# RMSprop (Root Mean Square Prop), Hinton et al. 2014

$$g_t = \nabla L(\theta_{t-1})$$

$$v_t = \gamma v_{t-1} + (1 - \gamma) g_t \odot g_t$$

$$\theta_t = \theta_{t-1} - \eta_t * \frac{g_t}{\sqrt{v_t + \epsilon}}$$

- $\theta_t, g_t, v_t$  is a vector;  $\eta_t, \gamma$  is a number
  - Hinton suggests  $\gamma$  to be set to 0.9
- RMSprop is an extension of Adagrad that deals with its aggressively decaying learning rate
- Sparse gradient has a larger learning rate and moves faster

# Adam (adaptive moment estimation), Kingma & Ba 2015

$$g_t = \nabla L(\theta_{t-1})$$

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t \odot g_t$$

$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

$$\theta_t = \theta_{t-1} - \eta_t * \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

- $\theta_t, g_t, m_t, v_t$  is a vector;  $\eta_t, \beta_1, \beta_2$  is a number
  - The authors suggest  $\beta_1 = 0.9, \beta_2 = 0.999$
- It works well for NLP and reinforcement learning
- It fails to achieve the target accuracy for ImageNet