

Real World Planning and Acting

CS4246/CS5446

AI Planning and Decision Making

This lecture will be
recorded!



Please join:
pollev.com/anarayan



Admin: Final Test

- Week 11 Lecture
- Date: April 3, 2024
- Time: 6:30 pm
- Duration: 90 min
- Venue: TBD
- Mode: Sit-in; using Exemplify



Topics

- Heuristics for planning (RN 11.3)
 - Problem relaxation and admissible heuristics
 - Domain independent planning (RN 11.3.1)
 - State abstraction in planning (RN 11.3.2)
- Hierarchical planning (RN 11.4)
 - Hierarchical task networks and HTN Planning
 - High-level actions (RN 11.4.1)
 - Searching for primitive solutions (RN 11.4.2)
 - Searching for abstract solutions (RN 11.4.3)



Recap

- Characteristics of problems classical planning can solve
 - Discrete
 - Deterministic
 - Static
 - Fully observable
- Representing a planning problem – PDDL
 - Power to deal with set operations
- Planning as a Search & Satisfiability problem

Quiz

Quiz answer

Quiz answer

Exercise

- Question:

- If the goal is $P \wedge Q \wedge R$
- If Action a_1 has effect $P \wedge Q \wedge \neg R$, is a_1 a relevant action?

Solving Planning Problems

Heuristics: for classical planning
HTN: New planning model

- Planning Problem or Model
 - Appropriate abstraction of states, actions, effects, and goals (and costs and values)
- Planning Algorithm
 - Input: a problem
 - Output: a solution in the form of an action sequence
- Planning Solution
 - A plan or path from the initial state(s) to the goal state(s)
 - Any path; OR
 - An optimal path wrt to costs or values
 - A goal state that satisfies certain properties



Heuristics for Planning

Improving Efficiency

Heuristics for Planning

- Heuristic function
 - $h(s)$ estimates distance from a state s to the goal g
- Main idea
 - Find **admissible heuristic** for distance
 - A* or other heuristic search can then find **optimal** solutions

Quiz

Quiz answer

Quiz answer

Heuristics for Planning

- Heuristic function

- $h(s)$ estimates distance from a state s to the goal g

- Main idea

- Find **admissible heuristic** for distance
 - A* or other heuristic search can then find **optimal** solutions

- Approach

- Define a **relaxed problem** that is easier to solve
 - Exact cost of solution to easier problem is **heuristic** for original problem

Types of Heuristics

Real-world problems use a combination of domain-independent and domain-specific heuristics

- Definition and application
 - Facilitated by factored representation for states and action schemas
- Search problem
 - A graph with states as nodes and actions as edges
 - Find a path connecting initial state to goal state
- Domain independent heuristics –Why?
 - Two ways to relax the problem (make it strictly easier):
 - Add more edges – easier to find path
 - Group multiple nodes together – abstract with fewer states and easier to search
 - Prune away irrelevant branches of search tree



Heuristics of Adding Edges

- Ignore preconditions heuristic
 - Drop all preconditions from actions
- Ignore selected preconditions heuristic
 - Derive simpler measures of distance from goal
- Ignore delete list heuristic
 - Allow monotonic progress toward goal
- Caveats:
 - Finding solution to relaxed problem is still NP-hard
 - Trade-off in optimality or admissibility sometimes required
 - Expensive to calculate heuristics
 - Do not reduce state-space size

How to Define an Edge-Based Heuristic?

- Approach:
 - Relax actions by removing all preconditions and all effects except goal literals
 - Count minimum no. of actions required for union effects to satisfy the goal
- Note:
 - An instance of the set cover problem – NP-hard
- Potential solution:
 - Simple greedy algorithm guaranteed to return a set covering whose size is within a factor of $\log(n)$ of the true minimum covering
 - where n is the number of literals in the goal
 - Loses guarantee of admissibility

Example: 8-Puzzle as Planning

- Planning as path search

- Game objective:

- To rearrange a given initial configuration (state) of eight numbered tiles arranged on a 3×3 board into given final or goal configuration (state)

Initial

2	■	3
1	8	4
7	6	5



Goal

1	2	3
8	■	4
7	6	5

Action(Slide(t, s_1, s_2))

Precond: $On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

Effect: $On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$)

Example: 8-Puzzle as Planning

- States

- A state description specifies the location of each of the eight tiles in one of the nine squares

- Actions or operators

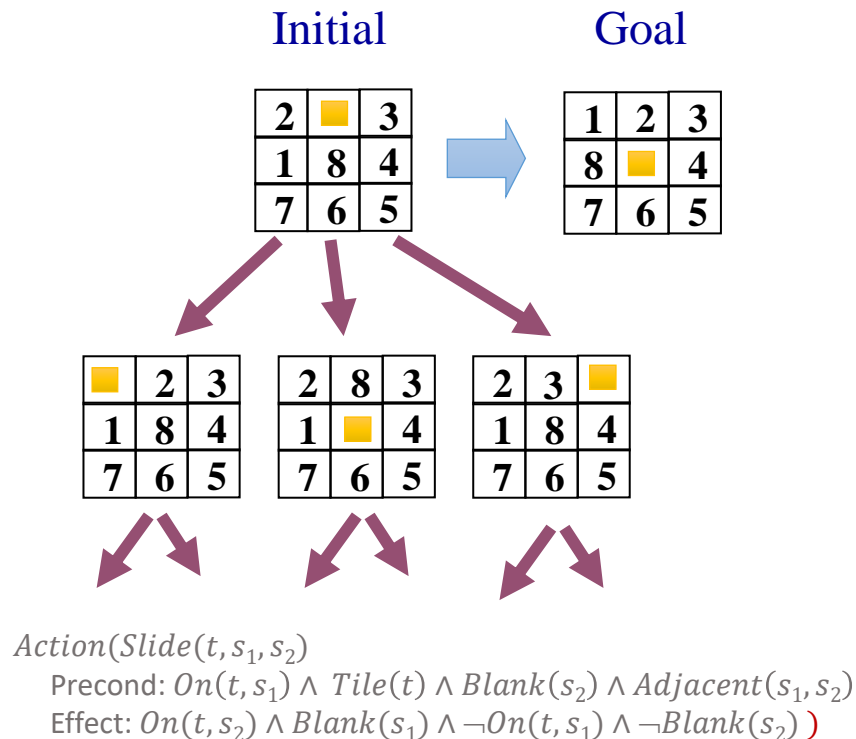
- Tile **slides** left, right, up, or down

- Goal test

- State matches the goal configuration

- Path cost

- Each step cost 1, so path cost is just the length of the path




Ignore Preconditions Heuristic

- Main idea:
 - All actions become applicable in all states
 - Any single goal fluent can be achieved in one step (if there is an applicable action)
 - No. of steps required to solve relaxed problem \approx no. of unsatisfied goals
 1. Some actions may achieve multiple goals
 2. Some actions may undo the effects of others
 - Possible accurate heuristic – consider 1 and ignore 2

Ignore Selected Preconditions Heuristic

- Remove Preconditions
 - $Blank(s_2) \wedge Adjacent(s_1, s_2)$
- What does the heuristic do?
 - No. of misplaced tiles heuristic
- What are the estimates?




$Action(Slide(t, s_1, s_2))$

Precond: $On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

Effect: $On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$

Ignore Selected Preconditions Heuristic

- Remove Preconditions
 - $Blank(s_2)$
- What does the heuristic do?
 - Manhattan-distance heuristic
- What are the estimates?
- Caveat:
 - Unclear which preconditions can be selectively ignored in general



$Action(Slide(t, s_1, s_2))$

Precond: $On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

Effect: $On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$



Ignore Delete List Heuristic

- Main idea:
 - Assume only positive literals in all plans and actions
 - Remove delete list from all actions (all negative literals from effects) to make monotonic progression toward goal
 - Create a relaxed version of original problem that is easier to solve, where solution length will serve as good heuristics
 - Approximate solution can be found in polynomial time by hill-climbing

Ignore Delete Lists Heuristic

- Ignore delete lists

- How does the problem become easier when ignore delete lists heuristic is applied to this schema?
- Goals are never undone
- Blanks always increase

Initial

2	■	3
1	8	4
7	6	5



Goal

1	2	3
8	■	4
7	6	5

Action(Slide(t, s_1, s_2))

Precond: $On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

Effect: $On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$

Can We Do Better?

- What would **YOU** do?
 - Can planning system emulate human moves?
 - Would that help?

Initial

2	■	3
1	8	4
7	6	5



Goal

1	2	3
8	■	4
7	6	5

Action(Slide(t, s_1, s_2))

Precond: $On(t, s_1) \wedge Tile(t) \wedge Blank(s_2) \wedge Adjacent(s_1, s_2)$

Effect: $On(t, s_2) \wedge Blank(s_1) \wedge \neg On(t, s_1) \wedge \neg Blank(s_2)$

Quiz

Quiz answer

Quiz answer

Domain-independent Pruning

- Symmetry reduction

- Prune all symmetric branches of the search tree except for one
- For many domains, efficiently solve intractable problems

How do you find preferred action(s)?

- Forward pruning

- Accept risk of pruning away an optimal solution, in order to focus search on promising branches

- Rule-out negative interactions

- A problem has **serializable subgoals** if there exists an order of subgoals such that the planner can achieve them in that order without having to undo any of the previously achieved subgoals



Serializable Subgoals

- Planning examples in the real world:
 - Build a tower of blocks on Table, in any order
 - Switch on all n lights with independent switches, in any order
 - Remote Agent Planner that commands NASA's Deep Space One spacecraft (1998) – serializable by design
 - able to command spacecraft in real-time



Heuristics of State Abstraction


- State abstraction:
 - Many-to-one mapping from states in the ground representation of the problem to the abstract representation
- Main approach:
 - Ignore some fluents
 - Solution in abstract state space will be shorter than a solution in the original space (assumes **admissible heuristic**)
 - Abstract solution extensible to solution for original problem

Example: Air Cargo Transportation



- **Original problem:**
 - 10 airports, 50 planes and 200 cargos
 - Total no. of states =
- **Assumption:**
 - All cargos are just in 5 airports, instead of 10, and all cargos in the same airport have the same destination.
- **Reformulation:**
 - Drop all the irrelevant At fluents (What are the relevant ones?)
 - Total no. of states =
- **Solution:**
 - Shorter than that for original problem (admissible heuristic)
 - Extensible by adding relevant actions

Decomposition



- Main idea:
 - Divide problem into parts, solving each part independently, and then combining the parts - Key idea in defining heuristics
- How to choose the right abstraction to reduce total cost?
 - Defining abstraction, doing abstract search, mapping abstraction back to original problem
 - Can the cost be less than original planning cost?
 - Example:
 - Pattern databases – cost of creation amortized over many problem instances

Planning Cost with Abstraction

- Problem definition:
 - Suppose the goal is a set of fluents G , divided into disjoint subsets G_1, \dots, G_n .
 - Find plans P_1, \dots, P_n that solve the respective subgoals
 - What is the estimated cost of the plan for achieving all of G ?
- Heuristic estimation:
 - Think of each $Cost(P_i)$ as a heuristic estimate
 - If each subproblem uses an admissible heuristic, taking Max is admissible
- Assuming subgoal independence:
 - Sum the cost of solving each independent subgoal; if admissible, Sum is better than Max – Why?
 - Solution **optimistic** when there are negative interactions between subplans for each subgoal; e.g., action in one subplan deletes goals in another subproblem.
 - Solution **pessimistic** (not admissible) when there are positive interactions; e.g., actions in one subplan achieves goals in another subproblem

Example: FF- FastForward

- Characteristics:

- Forward state-space searcher making use of effective heuristics
- Ignore-delete-list heuristic with graph plan for heuristic estimation
- Hill-climbing search (modified to keep track of plan) with heuristic to find a solution
 - Non-standard hill-climbing algorithm: avoids local maxima by running a breadth-first search from the current state until a better one is found
 - If this fails, FF switches to greedy best-first search instead
- [Hoffmann, 2001]

Classical Planning Today

Hint! Hint! Think about the project 😊

- Classical Planning research: Examples:
 - Families of systems in use – Heuristic Search Planner (HSP), Fast Forward (FF), Fast Downward
 - <https://planning.wiki/ref/planners/>
- International Planning Competitions
 - <https://www.icaps-conference.org/competitions/>
 - FastForward (FF) [Hoffmann, 2001] was the winner in the 2002 International Planning Competition (IPC)
 - SATPlan was the winner in the 2004 and 2006 IPC
 - IPC 2014 was won by SymBA*, based on bidirectional search using heuristics and abstraction.
 - IPC 2018 was won by Delfi, using deep learning to select a planner from a collection of planners including Fast Downward (uses forward search) with various heuristics and SymBA*
 - IPC 2020 – Hierarchical planning!
- Classical Planning in practice: Examples
 - Commercial video games [Neufeld, X., et al., 2019]
 - AI planning for enterprise [Sohrabi, S., 2019]
 - Space exploration [Estlin, T., et al. 07] [Rabideau, G., et al., 2020]



Summary

- Classical or deterministic planning
 - No consensus on the best approach; competition and cross-fertilization induce progress
- Using domain-independent heuristics
 - Transform planning problems into relaxed problem spaces
 - Effective heuristics derived with subgoal independence assumptions by:
 - relaxation of planning problem
 - pruning repeated or irrelevant branches
 - Solve with efficient algorithms



Homework

- Readings:

- Heuristics: RN: 11.3
- Hierarchical: RN: 11.4

- Reviews:

- RN: 3.3, 3.5, 3.6 (Review of search and heuristics in problem-solving)

References

- Content:
 - Nau, Dana. Lecture Slides from Automate Planning: Licensed under the Creative Commons Attribution-NonCommercial-ShareAlike License: <http://creativecommons.org/licenses/by-nc-sa/2.0/>
- Classical planning systems: (Journal articles publicly available online or through NUS Library e-Resources)
 - Hoffmann, J. and B. Nebel, *The FF planning system: fast plan generation through heuristic search*. J. Artif. Int. Res., 2001. **14**(1): p. 253–302.
 - Helmert, M., *The fast downward planning system*. J. Artif. Int. Res., 2006. **26**(1): p. 191–246.
 - Georgievski, I. and M. Aiello, *HTN planning: Overview, comparison, and beyond*. Artificial Intelligence, 2015. **222**: p. 124-156.
- Classical planning in the Real World:
 - Neufeld, X., et al., Building a Planner: A Survey of Planning Systems Used in Commercial Video Games. IEEE Transactions on Games, 2019. **11**(2): p. 91-108. <https://ieeexplore.ieee.org/document/8214211>
 - Sohrabi, S., AI Planning for Enterprise: Putting Theory Into Practice, in Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19. 2019, International Joint Conferences on Artificial Intelligence Organization. p. 6408--6410. <https://www.ijcai.org/proceedings/2019/0897.pdf>
 - Estlin, T., et al. *Increased Mars Rover Autonomy using AI Planning, Scheduling and Execution*. in *Proceedings 2007 IEEE International Conference on Robotics and Automation*. 2007. https://ai.jpl.nasa.gov/public/documents/papers/estlin_icra07_marsrover.pdf
 - Rabideau, G.; Wong, V.; Gaines, D.; Agrawal, J.; Chien, S.; Kuhn, S.; Fosse, E.; and Biehl, J. [Onboard Automated Scheduling for the Mars 2020 Rover](#). In *Proceedings of the International Symposium on Artificial Intelligence, Robotics and Automation for Space, of i-SAIRAS'2020*, Noordwijk, NL, 2020. European Space Agency https://ai.jpl.nasa.gov/public/documents/papers/M2020_OBP_i-SAIRAS2020_camera.pdf