

# CS6208 : Advanced Topics in Artificial Intelligence

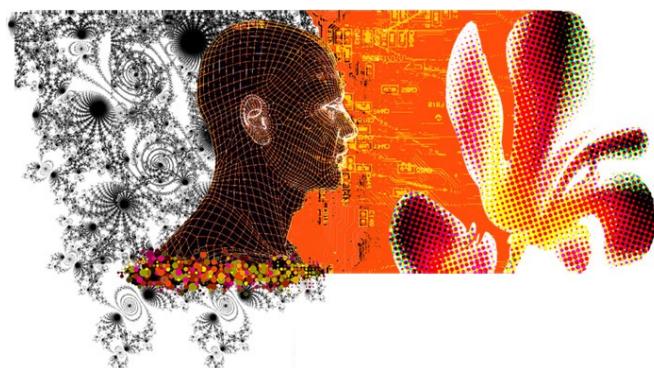
## Graph Machine Learning

### Lecture 5 : Graph Convolutional Networks Spectral and Spatial Techniques

Semester 2 2022/23

Xavier Bresson

<https://twitter.com/xbresson>



Department of Computer Science  
National University of Singapore (NUS)



# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

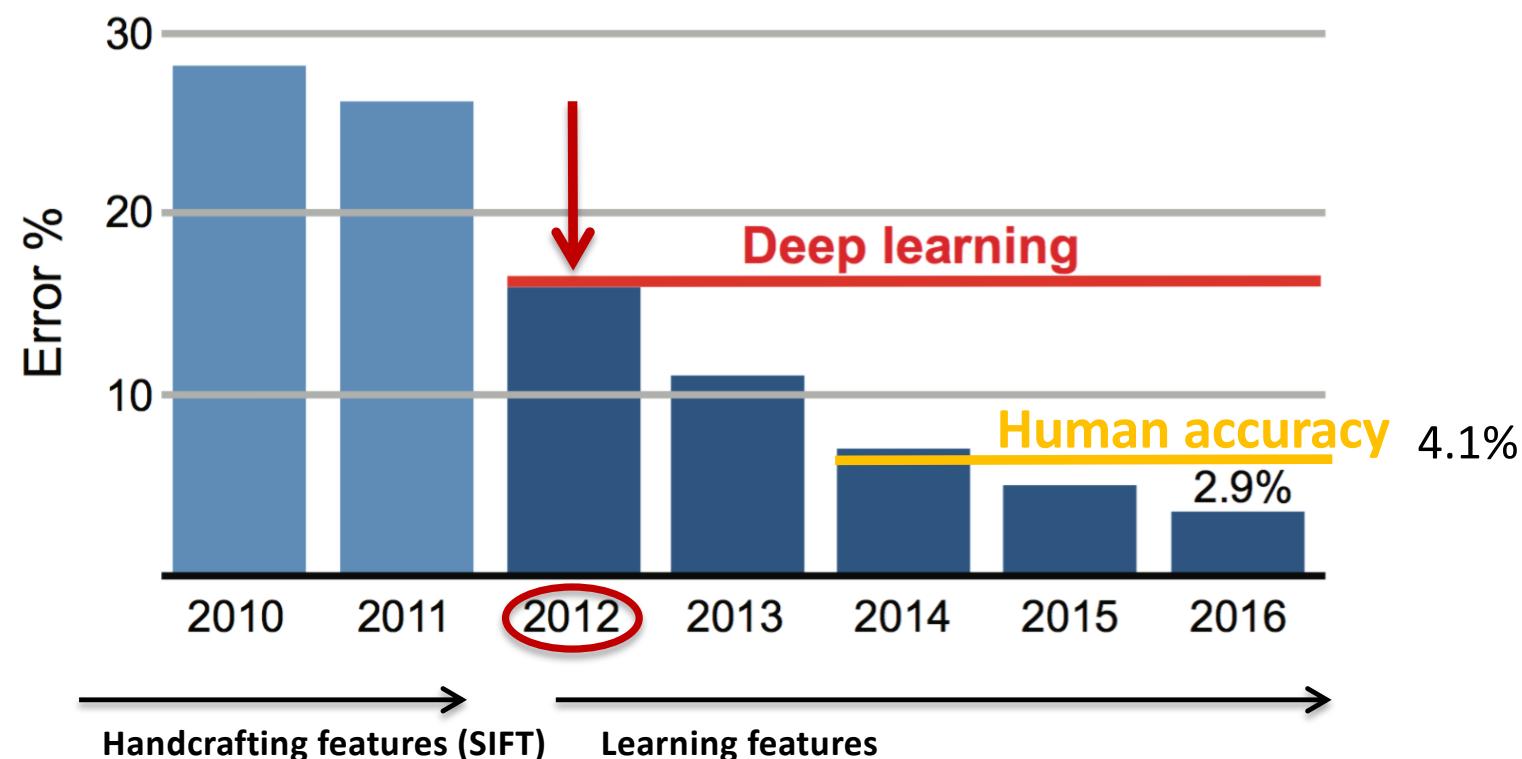
# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# ConvNets

- A breakthrough in Computer Vision :  
LeCun, Bottou, Bengio, Haffner 1998  
Krizhevsky, Sutskever, Hinton, 2012

IMAGENET



- Also in Speech and Natural Language Processing.

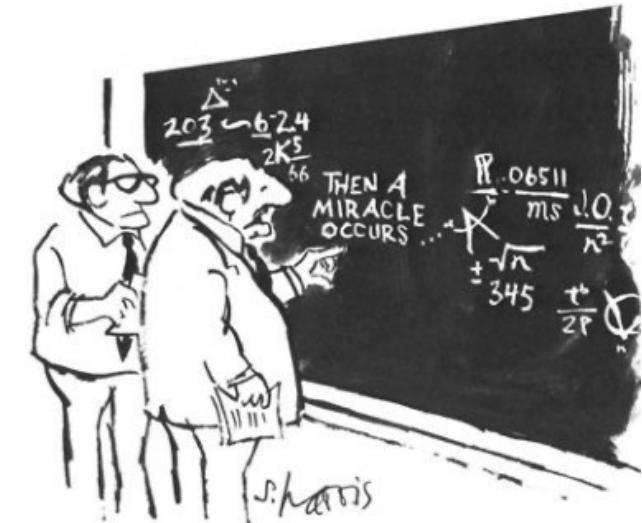
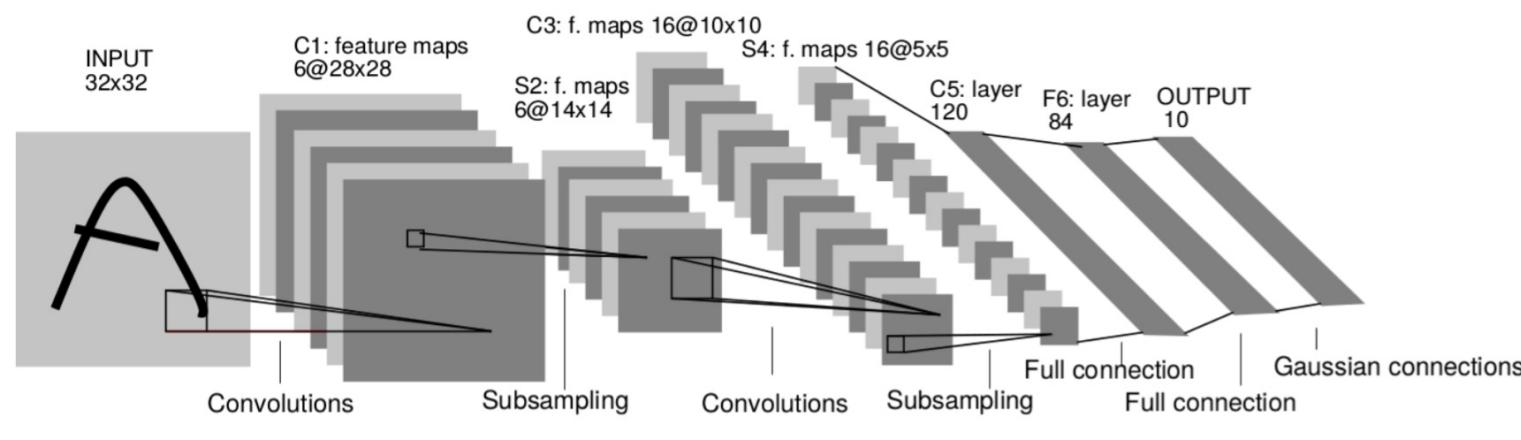


# ConvNets

- ConvNets are powerful architectures to solve high-dimensional learning problems.
- Curse of dimensionality :  
 $\text{dim(image)} = 1024 \times 1024 \approx 10^6$   
For  $N=10$  samples/dim  $\Rightarrow 10^{1,000,000}$  points

TECHNOLOGY The New York Times SUBSCRIBE NOW

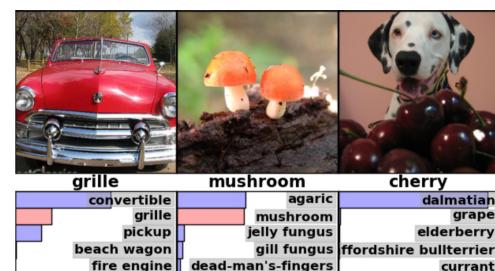
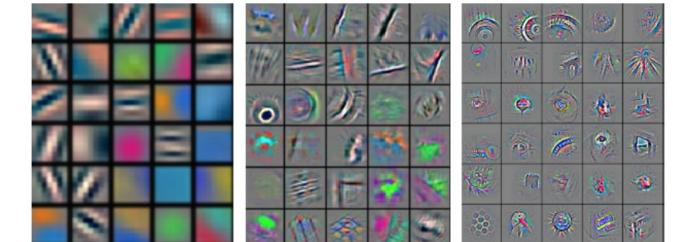
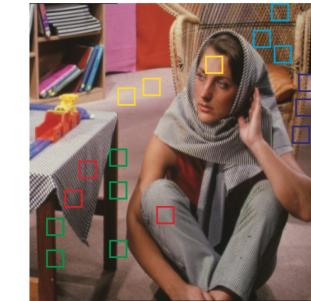
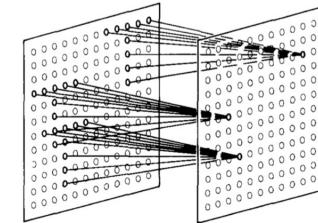
Turing Award Won by 3  
Pioneers in Artificial Intelligence



"I think you should be more explicit here in step two."

# ConvNets

- Main assumption :
  - Data (image, video, speech) is compositional, it is formed of patterns that are:
    - Local (Hubel-Wiesel 1962)
    - Stationary (shared patterns)
    - Hierarchical (multi-scale)
- ConvNets leverage the compositionality structure :
  - They extract compositional features and feed them to classifier, recommender, etc (end-to-end systems).



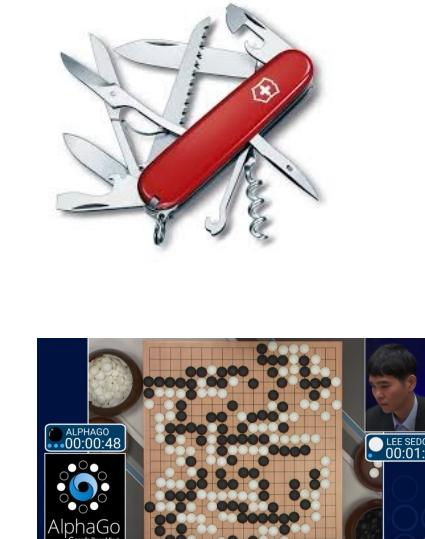
Computer Vision

```
/* Duplicate LSM field information. The lsm rule is opaque, so
 * re-initialized.
 */
static inline int audit_dupe_lsm_field(struct audit_field *df,
    struct audit_field *sf)
{
    int ret = 0;
    char lsm_str[1024];
    /* Our own copy of lsm_str */
    lsm_str = kstrdup(sf->lsm_str, GFP_KERNEL);
    if (unlikely(!lsm_str))
        return -ENOMEM;
    df->lsm_str = lsm_str;
    /* Our own (reflected) copy of lsm_rule */
    ret = security_audit_field_init(df->type, df->op, df->lsm_str,
        (void *)&df->lsm_rule);
    /* keep currently invalid fields around in case they
     * become valid after a policy reload.
     */
    if (ret == -EINVAL)
        pr_warn("audit rule for LSM '\\%s\\' is invalid\n",
            df->lsm_str);
    ret = 0;
}
return ret;
```

NLP



Speech



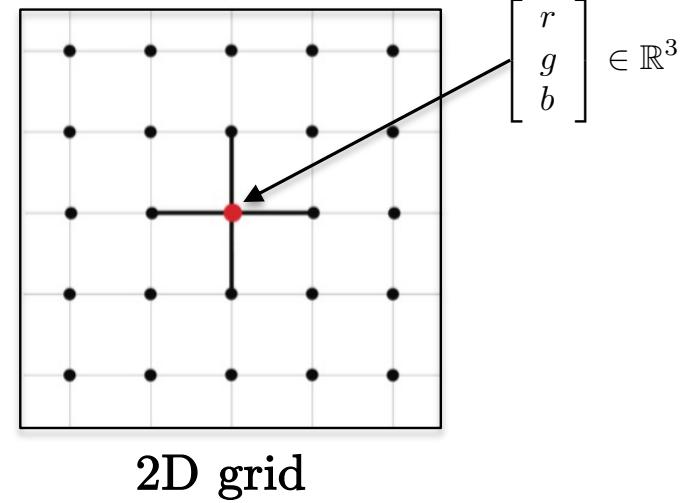
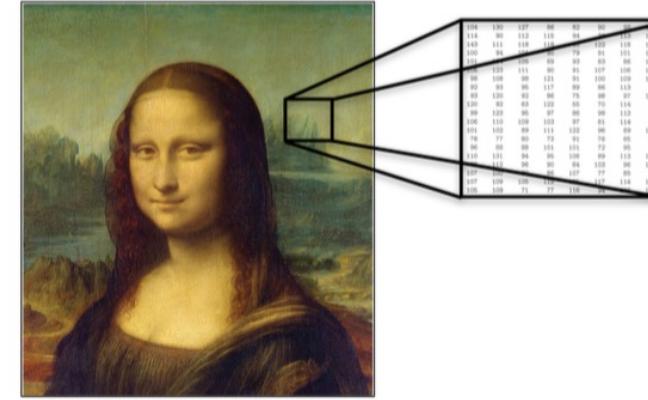
Game of Go

# Outline

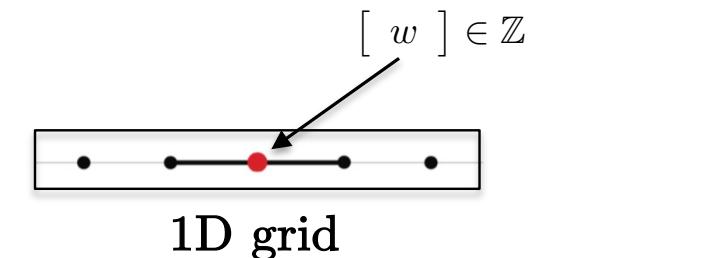
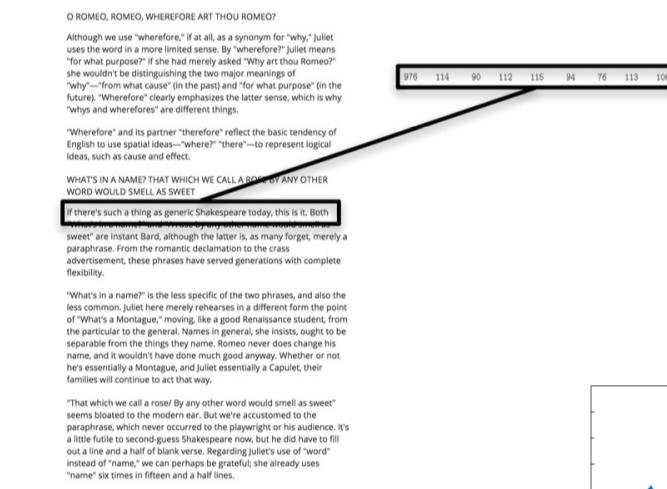
- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Data Domain

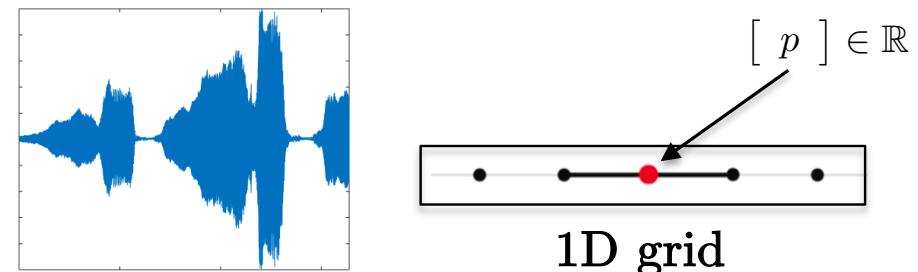
- Images, volumes, videos lie on  
2D, 3D, 2D+1 Euclidean domains (grids)



- Sentences, words, speech lie on  
1D Euclidean domain



- These domains have strong regular spatial structures.
  - All ConvNet operations are mathematically well defined and fast (convolution, pooling).



# Graph Domain



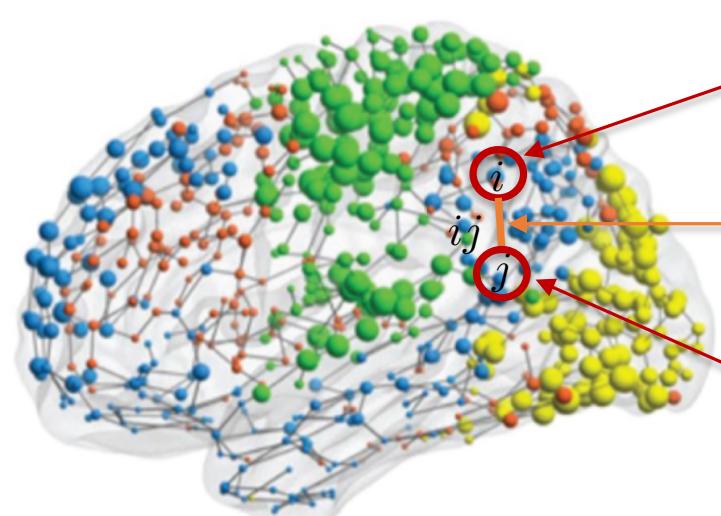
Social networks  
(Advertisement/  
recommendation)

Brain connectivity  
(sMRI)

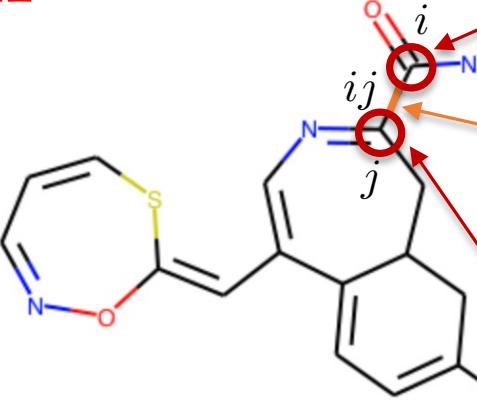
$$\text{User}_i \begin{bmatrix} \text{messages} \\ \text{images} \\ \text{videos} \end{bmatrix}_i \in \mathbb{R}^d$$

$$\text{User connection}_{ij} A_{ij} = \begin{cases} 1 & \text{if } ij \text{ friends} \\ 0 & \text{otherwise} \end{cases}$$

$$\text{User}_j \begin{bmatrix} \text{messages} \\ \text{images} \\ \text{videos} \end{bmatrix}_j \in \mathbb{R}^d$$



Brain analysis  
(Neuroscience/neuro-diseases)



$$\text{Atom}_i \begin{bmatrix} \text{type} \\ \text{coordinates} \\ \text{charge} \end{bmatrix}_i \in \mathbb{R}^{d_v}$$

$$\text{Bond}_{ij} A_{ij} = \begin{cases} 1 & \text{if } ij \text{ bond} \\ 0 & \text{otherwise} \end{cases}$$

$$\begin{bmatrix} \text{type} \\ \text{energy} \end{bmatrix}_{ij} \in \mathbb{R}^{d_e}$$

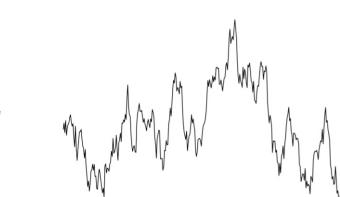
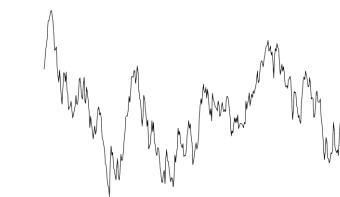
$$\text{Atom}_j \begin{bmatrix} \text{type} \\ \text{coordinates} \\ \text{charge} \end{bmatrix}_j \in \mathbb{R}^{d_v}$$

Quantum Chemistry  
(novel molecules for drugs  
and materials)

$$\text{ROI}_i \begin{bmatrix} a_1 \\ \vdots \\ a_T \end{bmatrix}_i \in \mathbb{R}^T$$

$$\text{Cerebral connection}_{ij} A_{ij} \in \mathbb{R}_+$$

$$\text{ROI}_j \begin{bmatrix} a_1 \\ \vdots \\ a_T \end{bmatrix}_j \in \mathbb{R}^T$$

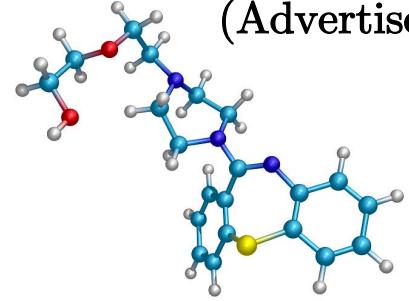


Functional  
activations (fMRI)

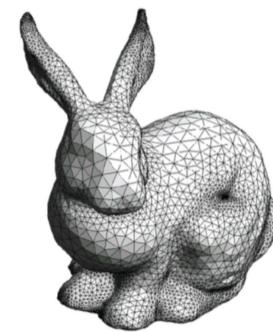
# Graph Domain



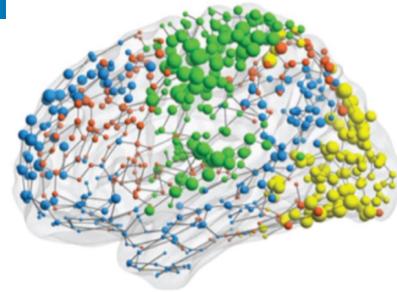
Social networks  
(Advertisement)



Drug/Material  
molecules  
(Chemistry)



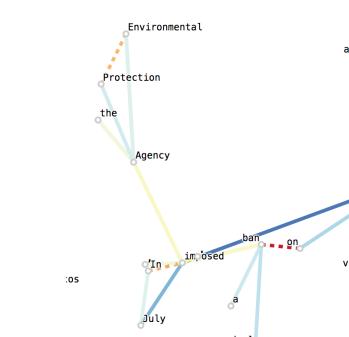
3D Meshes  
(Computer Graphics)



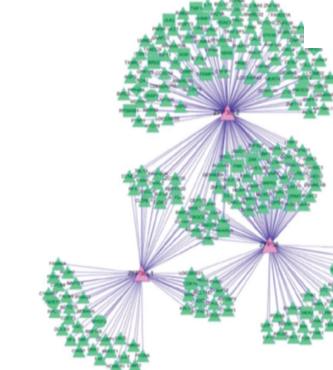
Brain  
connectivity  
(Neuroscience)



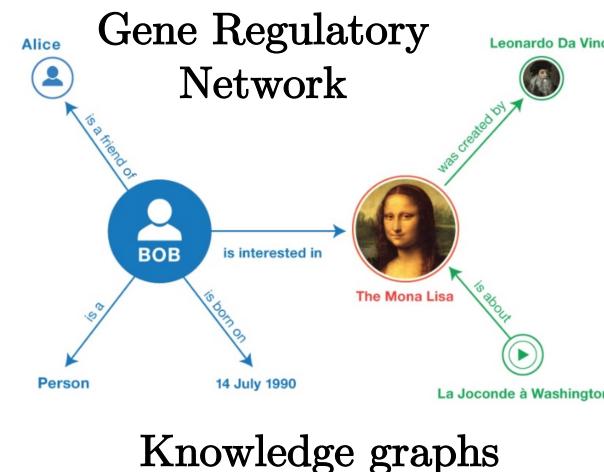
Transportation  
networks



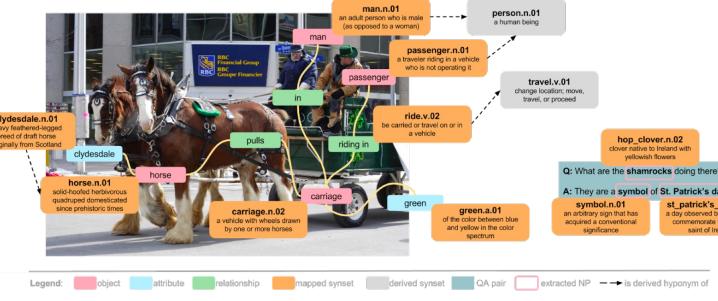
Words relationships  
(NLP)



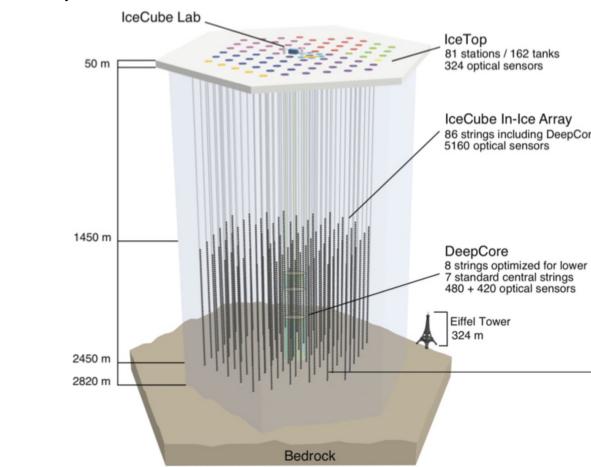
Recommender  
systems (Amazon,  
Netflix)



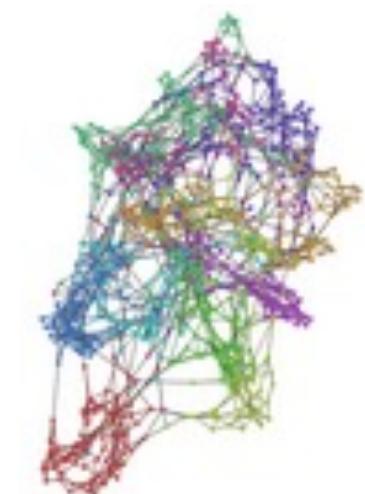
Knowledge graphs



Scene understanding



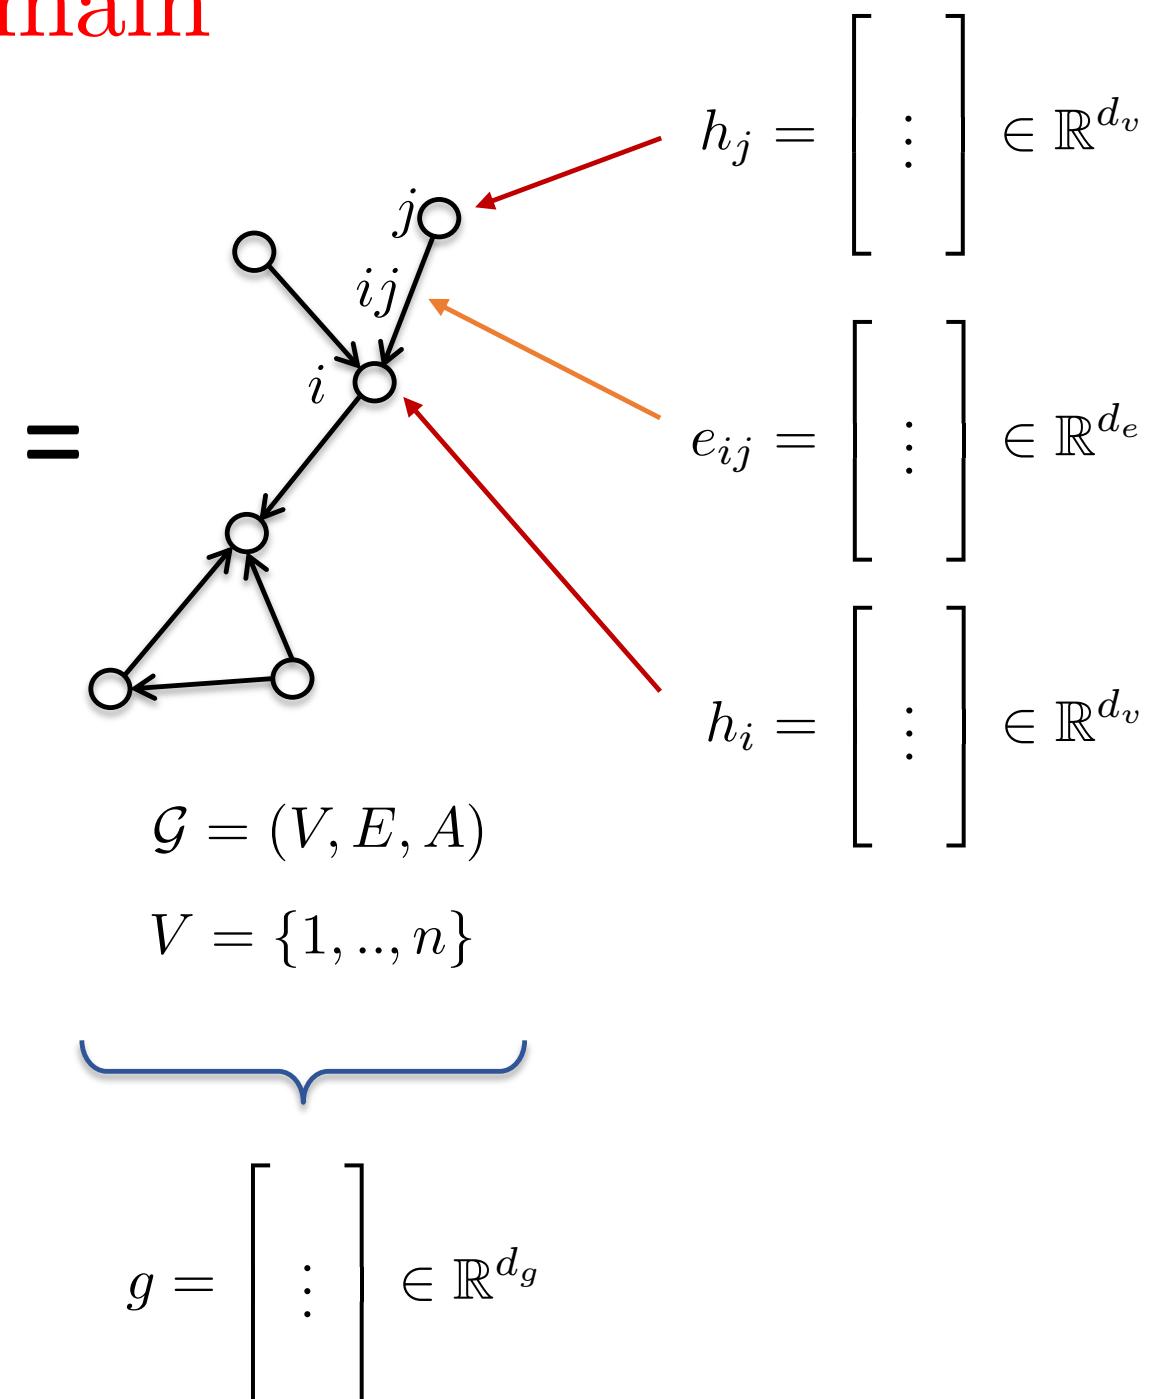
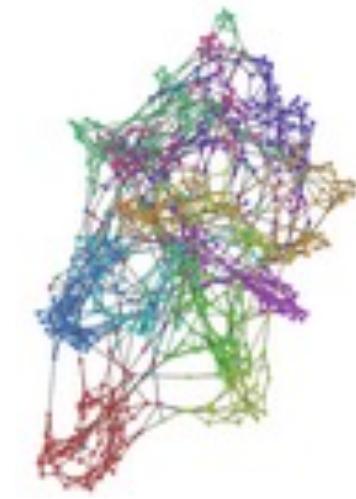
Neutrino  
detection (High-  
energy Physics)



Graph

# Graph Domain

- Graphs  $G$  are defined by :
  - Vertices  $V$
  - Edges  $E$
  - Adjacency matrix  $A$
- Graph features :
  - Node features :  $h_i, h_j$  (atom type)
  - Edge features :  $e_{ij}$  (bond type)
  - Graph features :  $g$  (molecule energy)



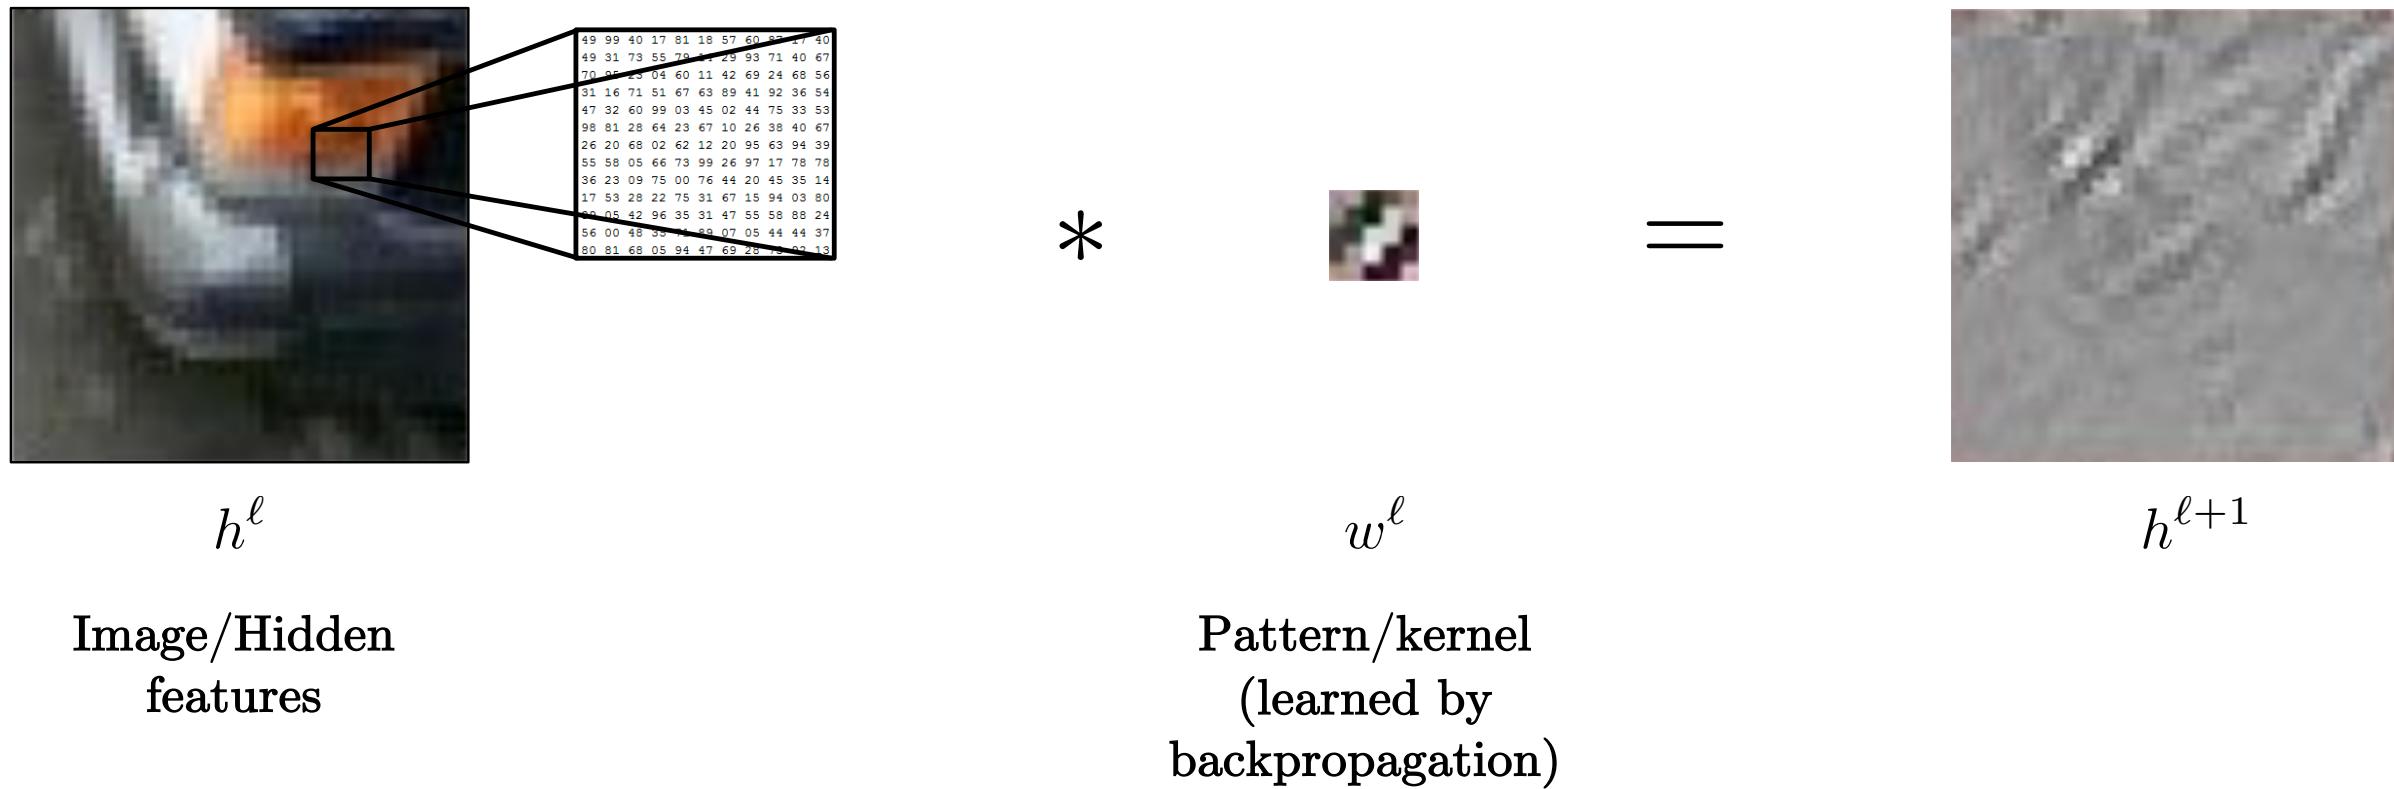
# Outline

- **Part 1 : Traditional ConvNets**
  - Architecture
  - Graph Domain
  - **Convolution**
- **Part 2: Spectral Graph ConvNets**
  - Spectral Convolution
  - Spectral GCNs
- **Part 3 : Spatial Graph ConvNets**
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Convolution

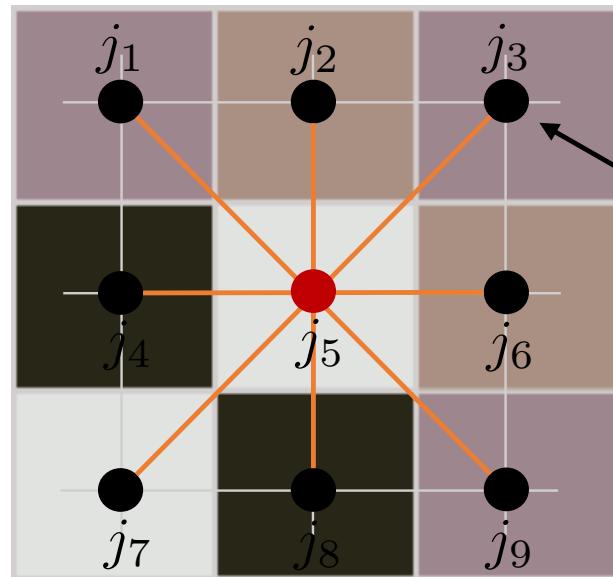
- Convolutional layer (for grids) :

$$h^{\ell+1} = w^\ell * h^\ell$$
$$n_1 \times n_2 \times d \quad n_1 \times n_2 \times d$$
$$3 \times 3 \times d$$



# Convolution

- How to define convolution ?
  - Definition #1 : Convolution as template matching
  - $O(n)$  by parallelization and for compact support patterns



All nodes of the template  $w^\ell$  are always ordered/positioned the same way !

$$w^\ell$$

Node  $j_3$  is always located at the top-right corner of the pattern.

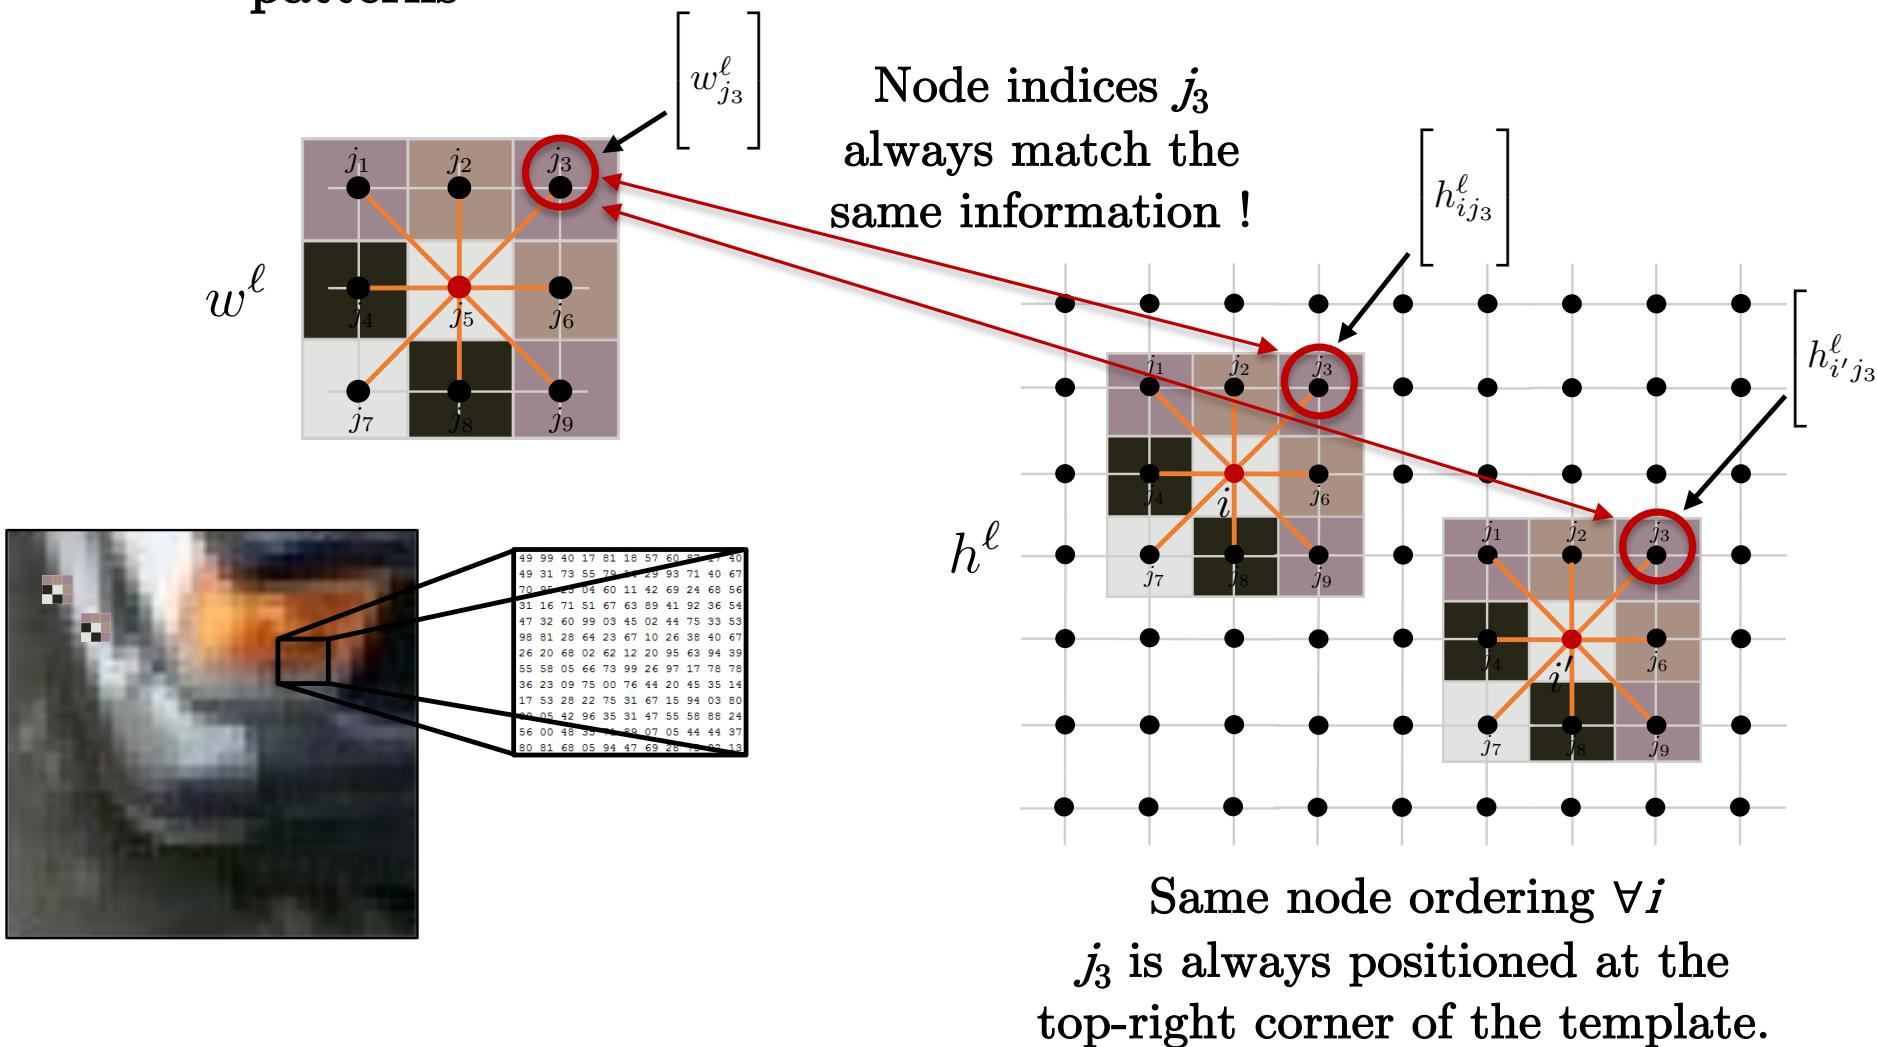
$$\begin{bmatrix} w_{j_3}^\ell \end{bmatrix} \in \mathbb{R}^d$$

Template features at  $j_3$

$$\begin{aligned}
 h_i^{\ell+1} &= w^\ell * h_i^\ell \\
 &= \sum_{\substack{j \in \Omega \\ j \in \mathcal{N}_i}} \langle w_j^\ell, \underbrace{h_{i-j}^\ell}_{h_{i+j}^\ell = h_{ij}^\ell} \rangle \\
 &= \sum_{j \in \mathcal{N}_i} \langle w_j^\ell, h_{ij}^\ell \rangle \\
 &= \sum_{j \in \mathcal{N}_i} \langle \begin{bmatrix} w_j^\ell \end{bmatrix}, \begin{bmatrix} h_{ij}^\ell \end{bmatrix} \rangle
 \end{aligned}$$

# Convolution

- How to define convolution ?
  - Definition #1 : Convolution as template matching
  - $O(n)$  by parallelization for compact support patterns

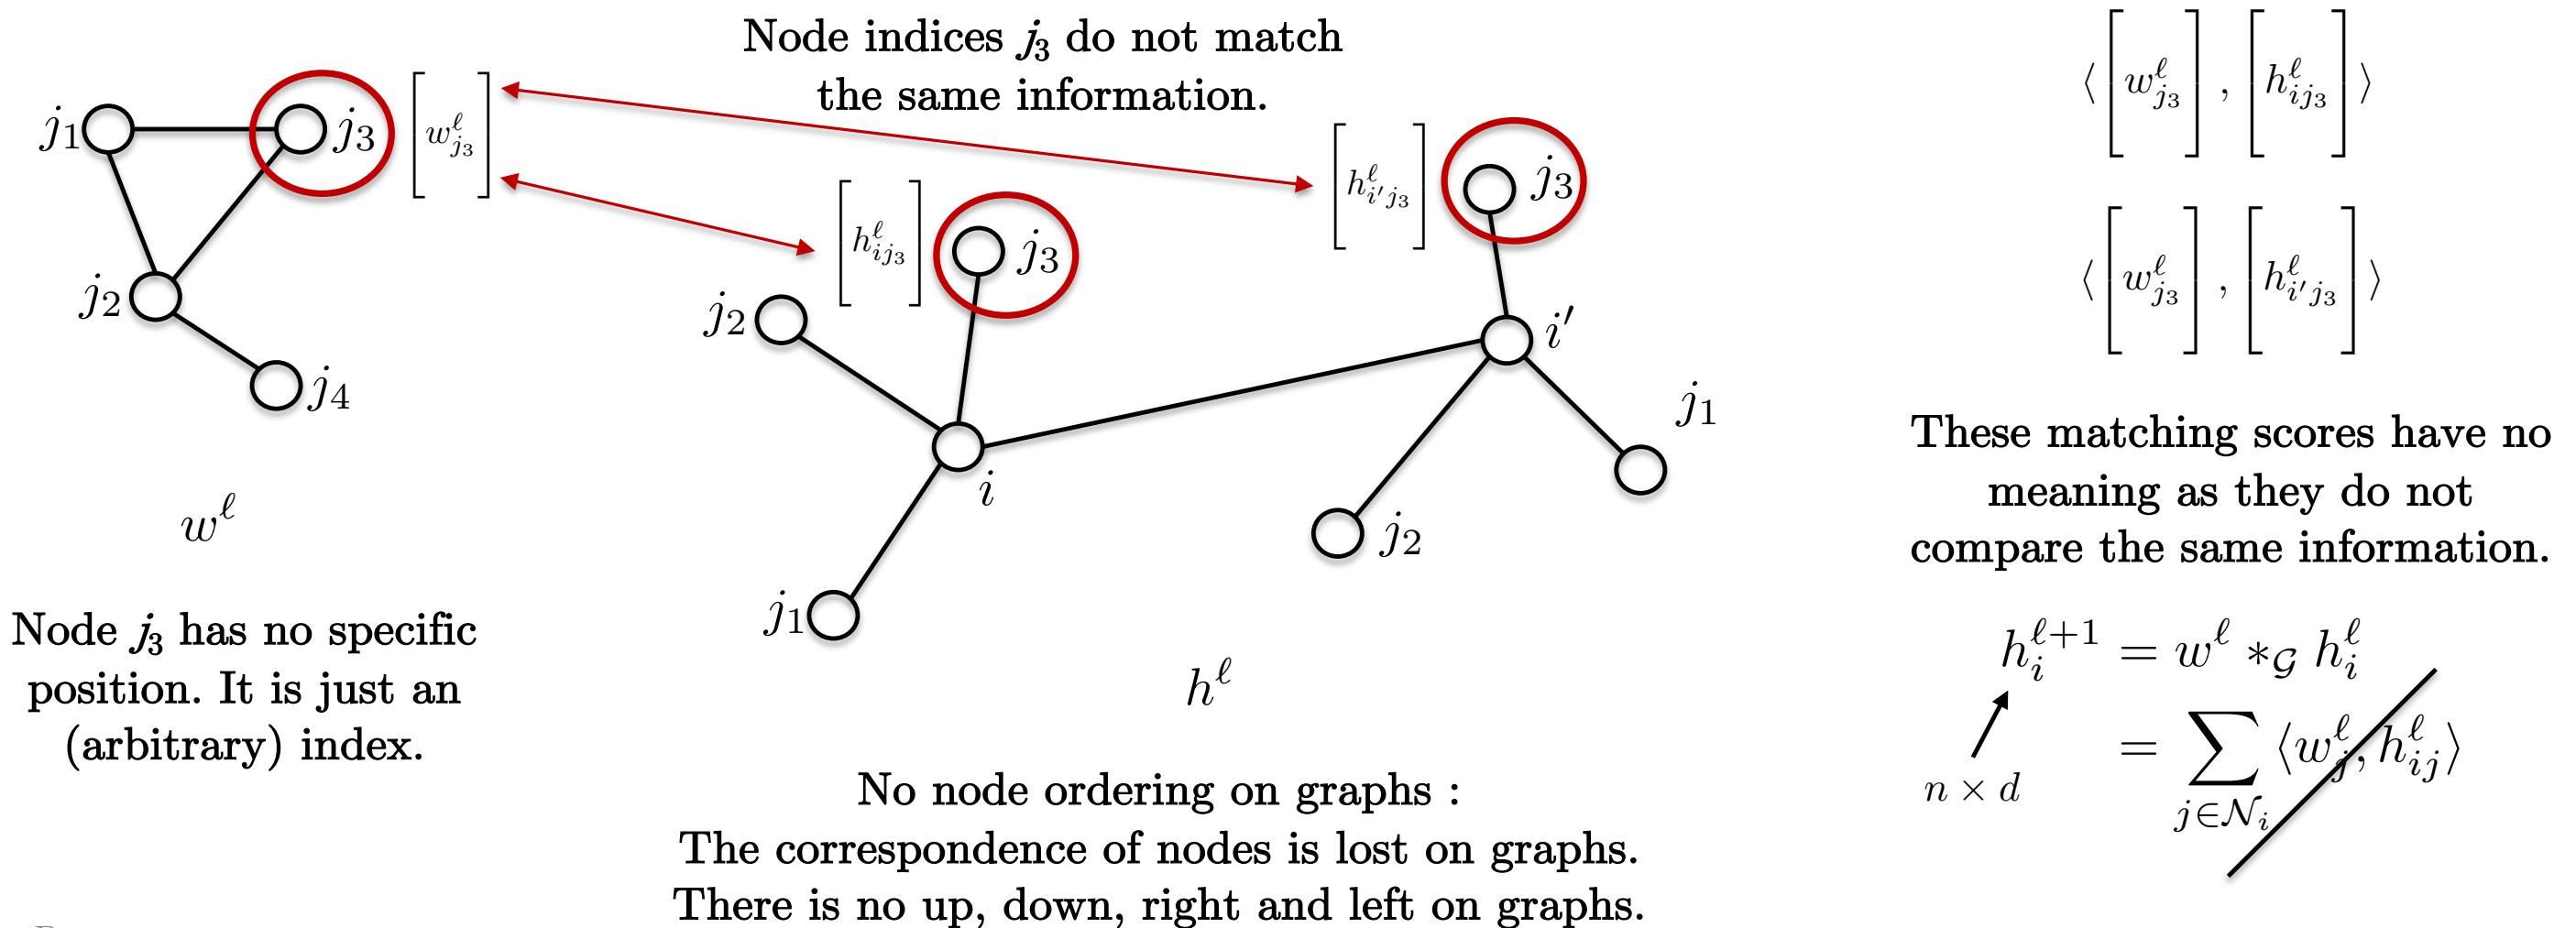


$$\begin{aligned}
 h_i^{\ell+1} &= w^\ell * h_i^\ell \\
 &= \sum_{j \in \mathcal{N}_i} \langle w_j^\ell, h_{ij}^\ell \rangle \\
 &= \sum_{j \in \mathcal{N}_i} \left\langle \begin{bmatrix} w_j^\ell \\ h_{ij}^\ell \end{bmatrix}, \begin{bmatrix} w_j^\ell \\ h_{ij}^\ell \end{bmatrix} \right\rangle \\
 &\quad \langle \begin{bmatrix} w_{j_3}^\ell \\ h_{ij_3}^\ell \end{bmatrix}, \begin{bmatrix} w_{j_3}^\ell \\ h_{ij_3}^\ell \end{bmatrix} \rangle \\
 &\quad \langle \begin{bmatrix} w_{j_3}^\ell \\ h_{i'j_3}^\ell \end{bmatrix}, \begin{bmatrix} w_{j_3}^\ell \\ h_{i'j_3}^\ell \end{bmatrix} \rangle
 \end{aligned}$$

These matching scores are always for the top-right corner between the template and the image patches.

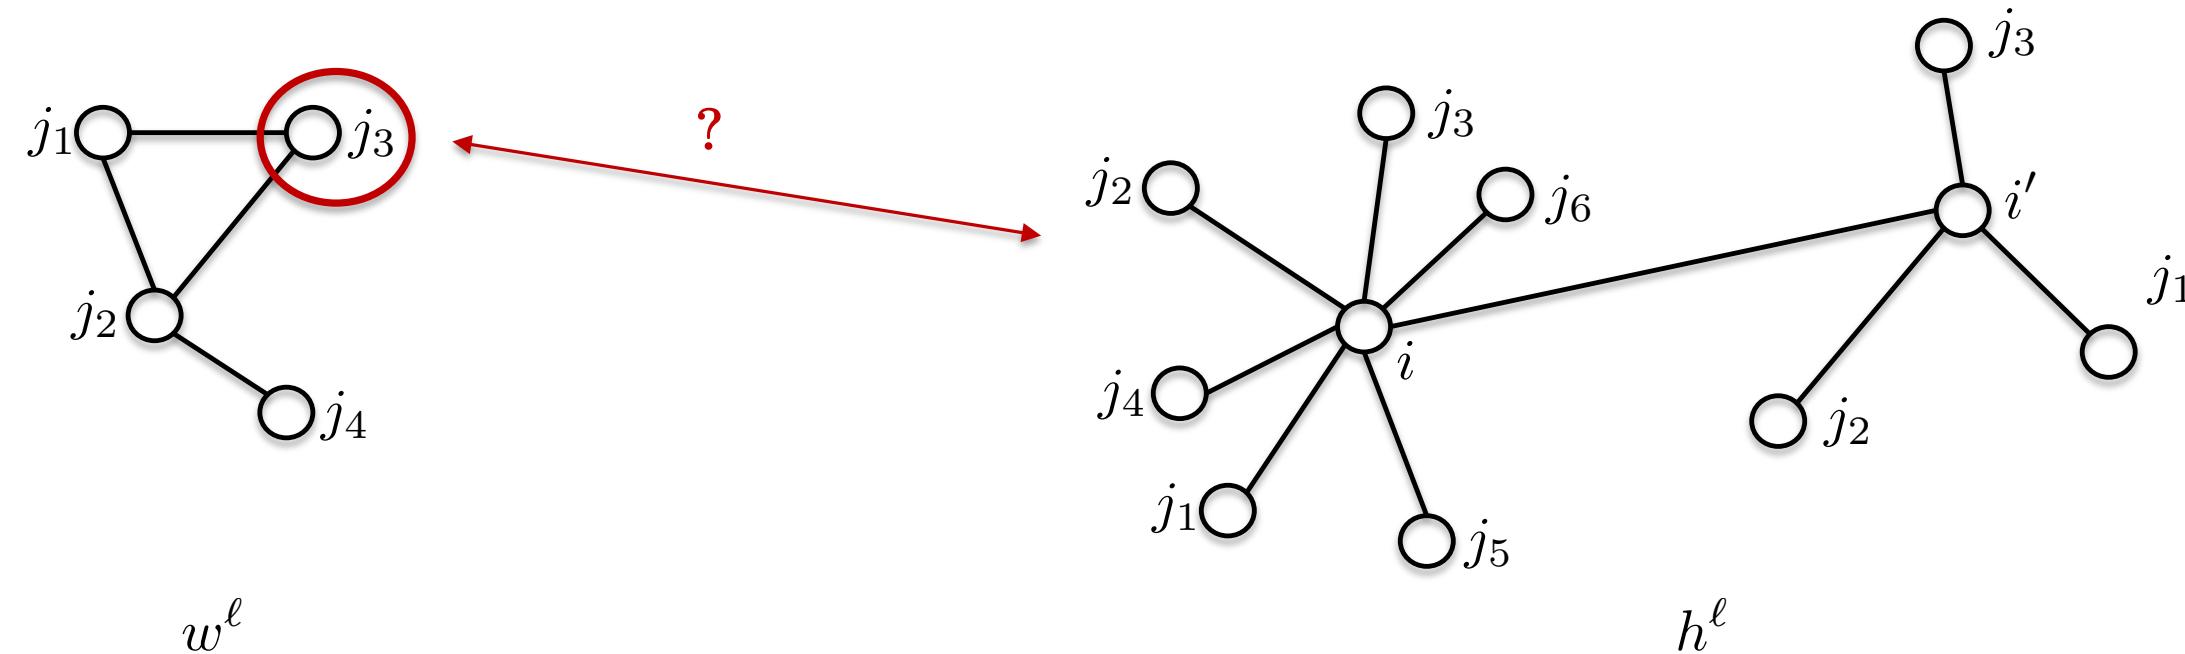
# Graph Convolution

- Can we extend template matching for graphs ?
- Main issues :
  - No node ordering : How to match template features with data features when nodes have no given position (index is not a position) ?



# Graph Convolution

- Can we extend template matching for graphs ?
  - Main issues :
    - No node ordering : How to match template features with data features ?
    - Heterogeneous neighborhood : How to deal with different neighborhood sizes ?



# Graph Convolution

- How to define convolution ?
  - Definition #1 : Template matching
  - Definition #2 : Convolution theorem
    - Fourier transform of the convolution of two functions is the pointwise product of their Fourier transforms
- Generic Fourier transform has  $O(n^2)$  complexity, but if the domain is a grid then complexity can be reduced to  $O(n \log n)$  with FFT<sup>[1]</sup>.
- Can we extend the Convolution theorem to graphs ?
  - How to define Fourier transform for graphs ?
  - How to compute fast spectral convolutions in  $O(n)$  time for compact kernels ?

$$w *_{\mathcal{G}} h \stackrel{?}{=} \mathcal{F}_{\mathcal{G}}^{-1}(\mathcal{F}_{\mathcal{G}}(w) \odot \mathcal{F}_{\mathcal{G}}(h))$$

[1] JW Cooley, JW Tukey, An algorithm for the machine calculation of complex Fourier series, 1965

# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Spectral Convolution

- Spectral graph theory
  - Book of Fan Chung<sup>[1]</sup> (harmonic analysis, graph theory, combinatorial problems, optimization)
- How to perform spectral convolution ?
  - Graph Laplacian
  - Fourier functions
  - Fourier transform
  - Convolution theorem

Conference Board of the Mathematical Sciences  
**CBMS**  
Regional Conference Series in Mathematics  
Number 92

Spectral Graph Theory  
Fan R. K. Chung

Published for the  
Conference Board of the Mathematical Sciences  
by the  
American Mathematical Society  
Providence, Rhode Island  
with support from the  
National Science Foundation

[1] FRK Chung, Spectral graph theory, 1997

# Graph Laplacian

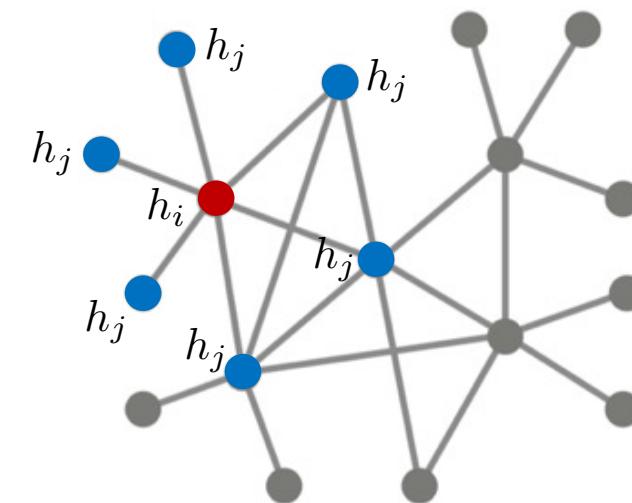
- Core operator in Spectral Graph Theory

$$\mathcal{G} = (V, E, A)_{n \times n} \rightarrow \Delta = I - D^{-1/2} A D^{-1/2} \quad \text{Normalized Laplacian}$$

where  $D = \text{diag}(\sum_{j \neq i} A_{ij})$

- Interpretation :
  - Measure of smoothness : Difference between local value  $h_i$  and its neighborhood average values  $h_j$ .

$$(\Delta h)_i = h_i - \sum_{j \in \mathcal{N}_i} \frac{1}{\sqrt{d_i d_j}} A_{ij} h_j$$



# Fourier Functions

- Eigen-decomposition of graph Laplacian :

Lap Eigenvalues/  
Spectrum

$$\Delta = \Phi^T \Lambda \Phi$$

$n \times n$

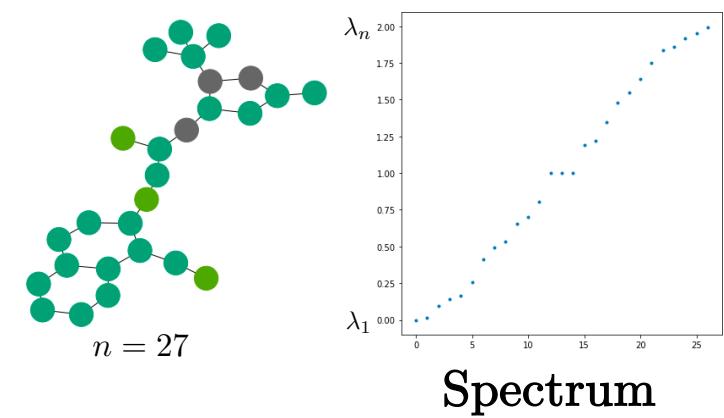
Lap Eigenvectors/  
Fourier functions

where  $\Phi = [\phi_1, \dots, \phi_n] = \begin{bmatrix} | & | \\ \phi_1 & \dots & \phi_n \\ | & | \end{bmatrix}$  and  $\Phi^T \Phi = I$ ,  $\langle \phi_k, \phi_{k'} \rangle = \delta_{kk'}$

where  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n) = \begin{bmatrix} \lambda_1 & 0 \\ & \ddots \\ 0 & \lambda_n \end{bmatrix}$  and  $0 \leq \lambda_1 \leq \dots \leq \lambda_n = \lambda_{\max} \leq 2$

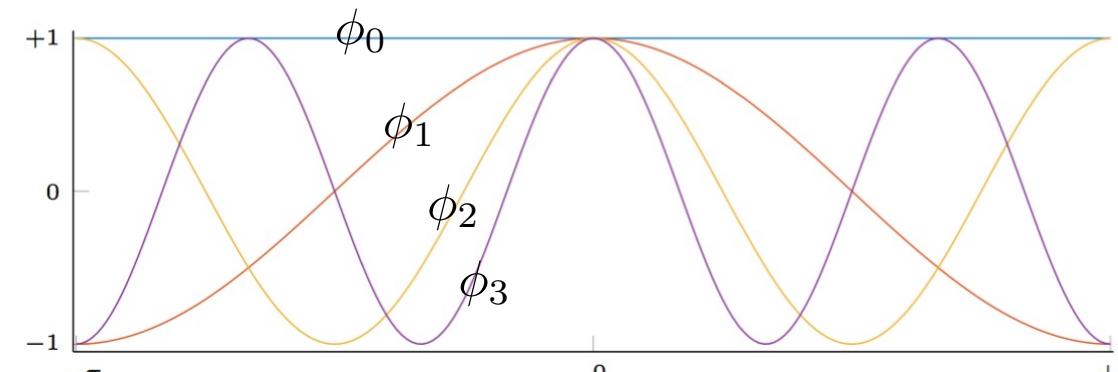
and  $\Delta \phi_k = \lambda_k \phi_k$ ,  $k = 1, \dots, n$

$n \times 1$



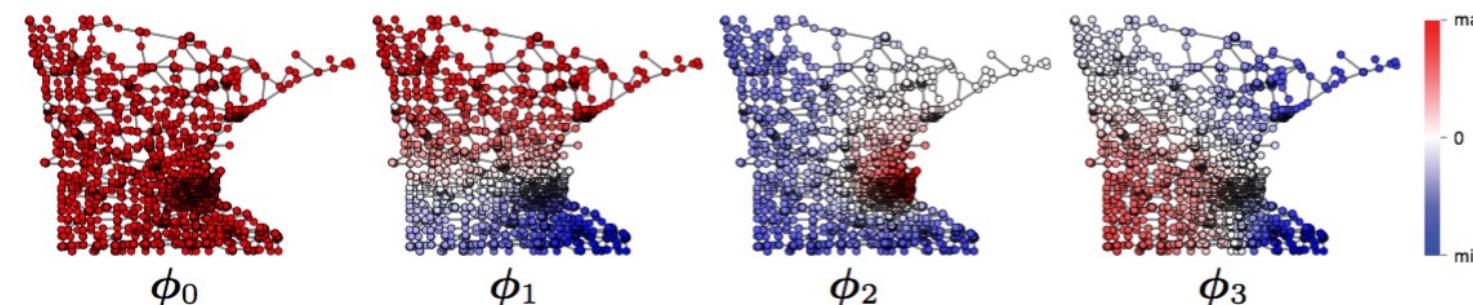
# Fourier Functions

- Grid/Euclidean domain :



First eigenvectors of 1D Euclidean Laplacian = standard Fourier basis

- Graph domain :



First Laplacian eigenvectors of a graph

Fourier functions related to graph geometry

(s.a. communities, hubs, etc)

Spectral graph clustering<sup>[1]</sup>

[1] Von Luxburg, A tutorial on spectral clustering, 2007

# Fourier Transform

- Fourier series : Decompose function  $h$  with Fourier functions<sup>[1]</sup> :

$$\begin{aligned}
 h &= \sum_{k=1}^n \underbrace{\langle \phi_k, h \rangle}_{\substack{\mathcal{F}(h)_k = \hat{h}_k = \phi_k^T h \\ \text{scalar}}} \phi_k \\
 &= \sum_{k=1}^n \hat{h}_k \phi_k \\
 &= \underbrace{\Phi \hat{h}}_{\mathcal{F}^{-1}(\hat{h})}
 \end{aligned}$$

Fourier transforms  
are one line of code  
(linear operations)

$$\left\{
 \begin{array}{ll}
 \mathcal{F}(h) = \Phi^T h & \text{Fourier Transform/} \\
 n \times 1 & \text{coefficients of Fourier Series} \\
 = \hat{h} & \\
 \mathcal{F}^{-1}(\hat{h}) = \Phi \hat{h} & \text{Inverse Fourier Transform} \\
 n \times 1 & \\
 = \Phi \Phi^T h = h & \text{as } \mathcal{F}^{-1} \circ \mathcal{F} = \Phi \Phi^T = I \quad \text{Orthonormal basis/} \\
 & \text{Invertible transformation}
 \end{array}
 \right.$$

[1] K. Hammond, P. Vandergheynst, R. Gribonval, Wavelets on graphs via spectral graph theory, 2011

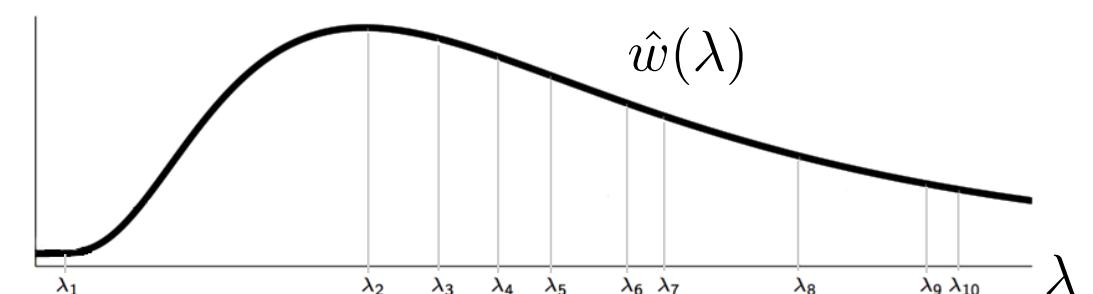
# Convolution Theorem

- Fourier transform of the convolution of two functions is the pointwise product of their Fourier transforms :

$$\begin{aligned}
 w * h &= \underbrace{\mathcal{F}^{-1}}_{\Phi} \left( \underbrace{\mathcal{F}(w)}_{\Phi^T w = \hat{w}} \odot \underbrace{\mathcal{F}(h)}_{\Phi^T h} \right) \\
 &= \underbrace{\Phi}_{n \times n} \left( \underbrace{\hat{w}}_{n \times 1} \odot \underbrace{\Phi^T h}_{n \times 1} \right) \\
 &= \Phi \left( \underbrace{\hat{w}(\Lambda)}_{n \times n} \underbrace{\Phi^T h}_{n \times 1} \right) \\
 &= \Phi \hat{w}(\Lambda) \Phi^T h \\
 &= \underbrace{\hat{w}(\Phi \Lambda \Phi^T)}_{\Delta} h \\
 &= \underbrace{\hat{w}(\Delta) h}_{n \times n \quad n \times 1}
 \end{aligned}$$

$$\hat{w} = \begin{bmatrix} \hat{w}(\lambda_1) \\ \vdots \\ \hat{w}(\lambda_n) \end{bmatrix}$$

$$\hat{w}(\Lambda) = \text{diag}(\hat{w}) = \begin{bmatrix} \hat{w}(\lambda_1) & & 0 \\ & \ddots & \\ 0 & & \hat{w}(\lambda_n) \end{bmatrix}$$



Expensive computation  $O(n^2)$   
No FFT

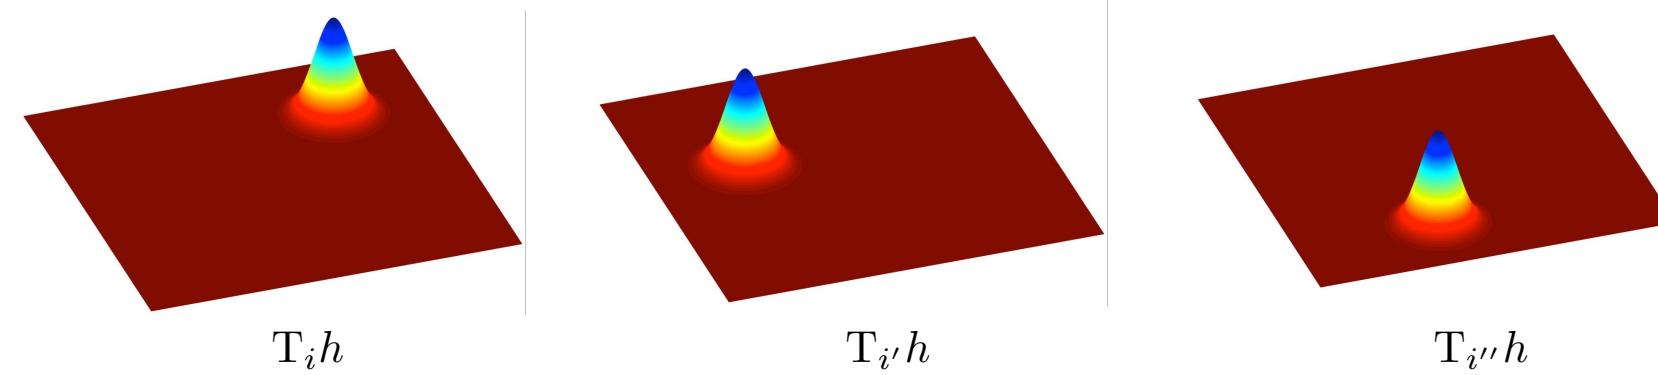
Spectral function/filter

# No Shift Invariance for Graphs

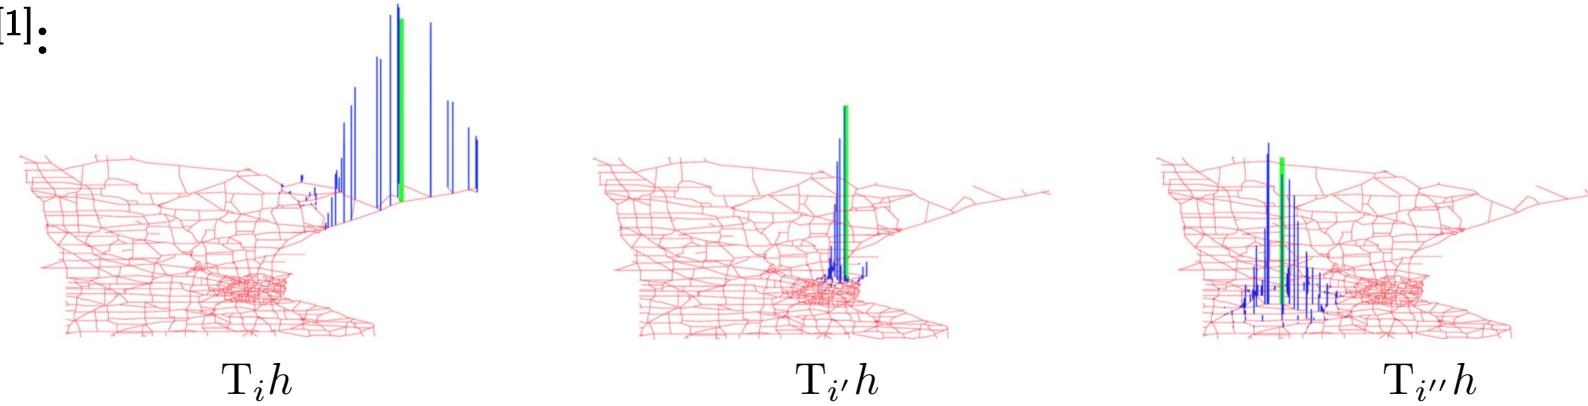
- A signal  $h$  on graph can be translated to vertex  $i$  as follows :

$$T_i h = \delta_i * h$$

- Grid/Euclidean domain :



- Graph domain<sup>[1]</sup>:



[1] K. Hammond, P. Vandergheynst, R. Gribonval, Wavelets on graphs via spectral graph theory, 2011

# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Vanilla Spectral GCN<sup>[1]</sup>

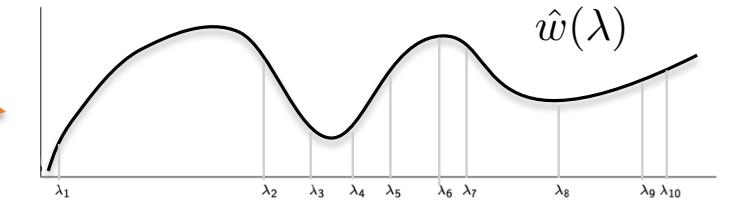
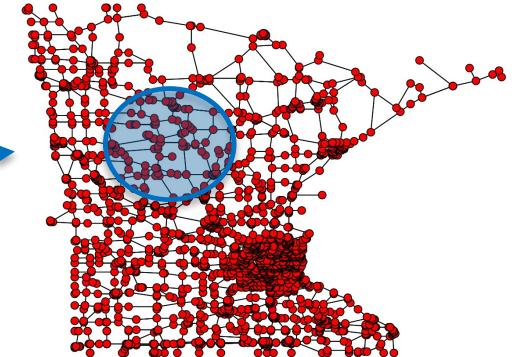
- Graph spectral convolutional layer :

$$\begin{aligned}
 h^{\ell+1} &= \eta(w^\ell * h^\ell) \\
 &= \eta(\hat{w}^\ell(\Delta)h^\ell) \\
 &= \eta(\Phi \hat{w}^\ell(\Lambda) \Phi^T h^\ell)
 \end{aligned}$$

$n \times d$        $n \times n$        $n \times d$

Spatial filter

Spectral filter



- First spectral technique for ConvNets
- Limitations :
  - No guarantee of spatial localization of filters
  - $O(n)$  parameters to learn per layer
  - $O(n^2)$  learning complexity (Fourier transform with full matrix  $\phi$ )

$$\hat{w}(\Lambda) = \begin{bmatrix} \hat{w}(\lambda_1) & 0 \\ 0 & \hat{w}(\lambda_n) \end{bmatrix}$$

[1] Bruna, Zaremba, Szlam, LeCun, Spectral Networks and Locally Connected Networks on Graphs, 2014

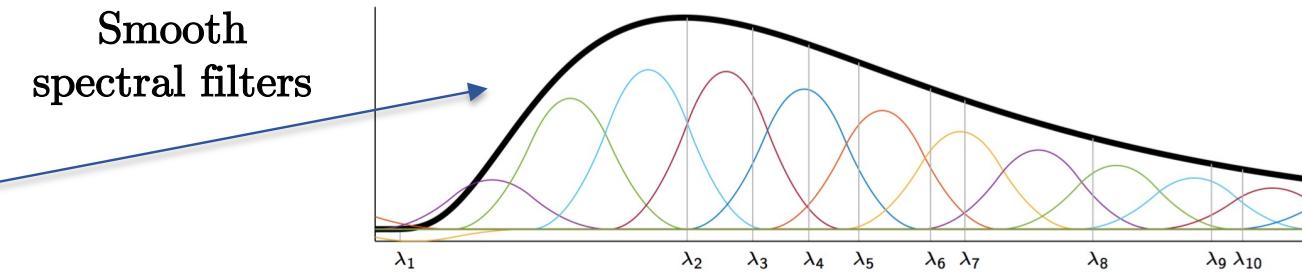
# SplineGCNs[1,2]

- Graph spectral convolutional layer :

$$\begin{aligned}
 h^{\ell+1} &= \eta( w^\ell * h^\ell ) \\
 &= \eta( \Phi \hat{w}^\ell(\Lambda) \Phi^T h^\ell ) \\
 \hat{w}^\ell(\Lambda) &= \text{diag}(B w^\ell) \\
 &\quad \xrightarrow{n \times n} K \times 1 \\
 &\quad \xrightarrow{n \times K} n \times n
 \end{aligned}$$

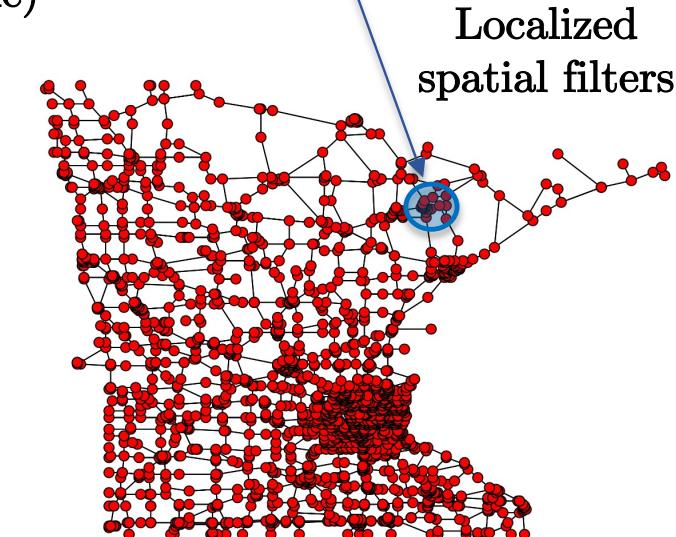
The  $K$  coefficients  $w^\ell$  are learned by backpropagation

Smooth spectral filters / Linear combination of  $K$  smooth kernels  $B$  (splines)



Parseval's Identity  
 Localization in space  $\Leftrightarrow$   
 Smoothness in frequency domain  
 (Heisenberg's uncertainty principle)

$$\int |x|^{2k} |w(x)|^2 dx = \int \left| \frac{\partial^k \hat{w}(\lambda)}{\partial \lambda^k} \right|^2 d\lambda$$



- Localized filters in space (fast-decaying)
- $O(1)$  parameters to learn per layer
- $O(n^2)$  learning complexity (Fourier transform with full matrix  $\phi$ )

[1] Bruna, Zaremba, Szlam, LeCun, Spectral Networks and Locally Connected Networks on Graphs, 2014  
 [2] Henaff, Bruna, LeCun, Deep Convolutional Networks on Graph-Structured Data, 2015

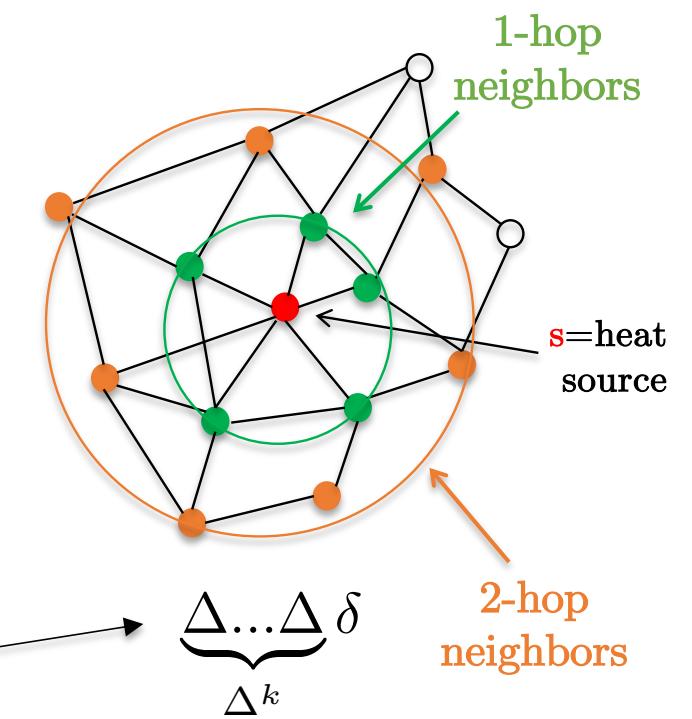
# LapGCNs<sup>[1]</sup>

- How to learn in linear time  $O(n)$  (w.r.t. graph size  $n$ ) ?
    - $O(n^2)$  complexity comes from the direct use of Laplacian eigenvectors :
- $$O(w * h) = O(\Phi \hat{w}^\ell(\Lambda) \Phi^T h) = O(n^2)$$
- $n \times d \quad n \times n \quad n \times n \quad n \times d$   
Full matrix
- How to avoid the eigen-decomposition ?
  - Learn directly functions of the Laplacian !

$$w * h = \hat{w}(\Delta)h$$

$$\hat{w}(\Delta) = \sum_{k=0}^{K-1} w_k \underbrace{\Delta \dots \Delta}_{\Delta^k} \delta$$

Coefficients  $w_k$  are learned by backpropagation.



Filters are exactly localized in  $k$ -hop supports.  
(each Laplacian operation increases the support of a function by 1 hop)

[1] Defferrard, Bresson, Vandergheynst, Convolutional neural networks on graphs with fast localized spectral filtering, 2016

# LapGCNs

- Learning complexity :

$$\begin{aligned} w * h &= \hat{w}(\Delta)h \\ &= \sum_{k=0}^{K-1} w_k \Delta^k h \\ &= \sum_{k=0}^{K-1} w_k X_k, \text{ with } X_k = \underset{n \times n}{\Delta} X_{k-1} \text{ and } X_0 = \underset{n \times d}{h} \end{aligned}$$

Recursive  
equation

- Sequence  $\{X_k\}$  is generated by multiplying a matrix  $\Delta$  and a vector  $X_{k-1}$   $\Rightarrow$  Complexity is  $O(E.K)=O(n)$  for sparse (real-world) graphs.
- No eigen-decomposition of Laplacian ( $\phi, \Lambda$ ) was required.
  - The name spectral GCNs can be misguided as the Lap eigen-decomposition is not used (computations are done in the spatial domain, not the spectral domain).
  - Graph convolutional layers are (sparse) linear operations, thus GPU friendly (but not yet optimized).

# LapGCNs

- Implementation :

$$\begin{aligned}
 h^{\ell+1} &= \eta\left(\sum_{k=0}^{K-1} w_k^\ell \Delta^k h^\ell\right) \\
 &= \eta\left(\sum_{k=0}^{K-1} w_k^\ell X_k\right) \\
 &= \eta\left((w^\ell)^T \bar{X}\right) \\
 &\quad \underbrace{\qquad\qquad\qquad}_{\text{reshape}} \\
 &\quad \qquad\qquad\qquad n \times d \\
 \text{with } \bar{X} &= \begin{bmatrix} - & \bar{X}_0 & - \\ & \vdots & \\ K \times nd & -\bar{X}_{K-1}- \end{bmatrix} \\
 \bar{X}_k &= \text{reshape}(X_k) \\
 1 \times nd & \qquad\qquad\qquad n \times d
 \end{aligned}$$

- Filters are exactly localized in  $K$ -hop support
- $O(1)$  parameters to learn per layer
- $O(n)$  learning complexity
- Monomials basis are unstable under coefficients perturbation (hard to optimize)

$$1, x, x^2, x^3, \dots \rightarrow \Delta^0, \Delta^1, \Delta^2, \Delta^3, \dots$$

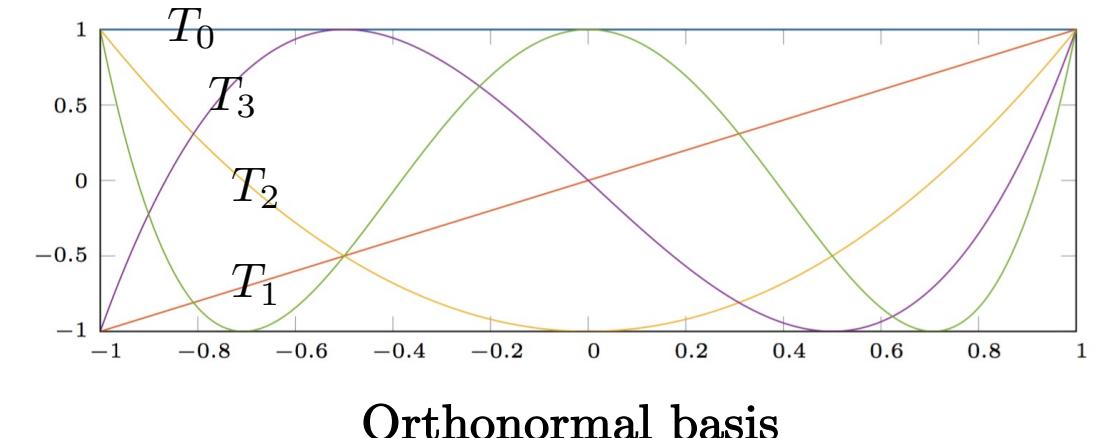
$$w^\ell = \begin{bmatrix} w_0^\ell \\ \vdots \\ w_{K-1}^\ell \end{bmatrix}$$

# Chebyshev Polynomials

- Graph spectral convolution with Chebyshev polynomials :

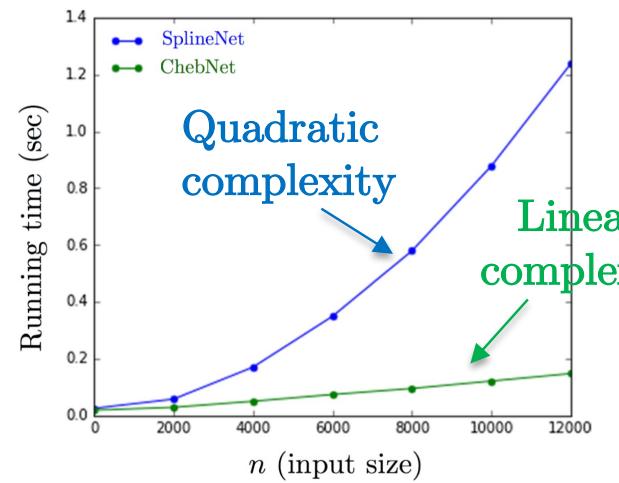
$$\begin{aligned}
 w * h &= \hat{w}(\Delta)h \\
 &= \sum_{k=0}^{K-1} w_k T_k(\Delta)h \\
 &= \sum_{k=0}^{K-1} w_k X_k, \text{ with } X_k = 2\tilde{\Delta}X_{k-1} - X_{k-2}, X_0 = h, X_1 = \tilde{\Delta}h \text{ and } \tilde{\Delta} = 2\lambda_n^{-1}\Delta - I
 \end{aligned}
 \quad \text{Recursive equation}$$

- Filters are exactly localized in  $K$ -hop support
- $O(1)$  parameters to learn per layer
- $O(n)$  learning complexity
- Stable under coefficients perturbation

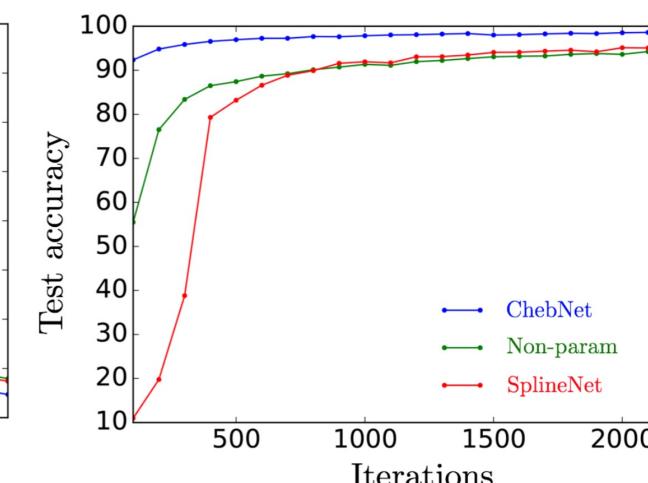
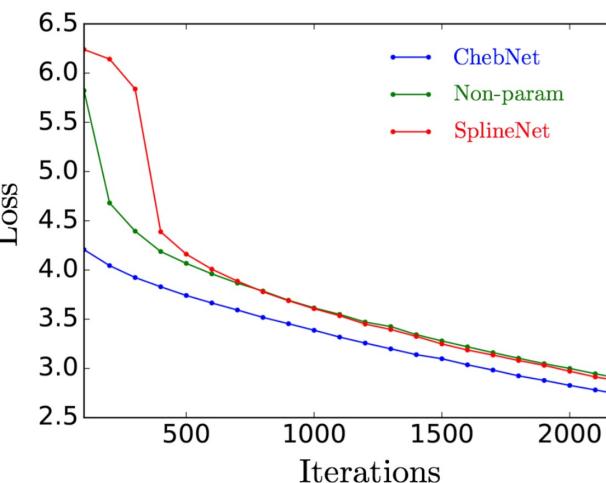


# MNIST Numerical Experiment<sup>[1]</sup>

Running time



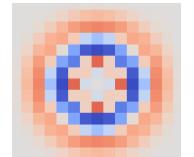
Optimization



Accuracy

Model	Order	Accuracy
LeNet5	-	99.33%
SplineNet	25	97.75%
ChebNet	25	99.14%

- ChebNets :
  - ConvNets for arbitrary graph domains
  - Same  $O(n)$  learning complexity (but larger complexity constant)
  - Limitation : Isotropic model
  - Isotropy vs anisotropy
    - Standard ConvNets produce anisotropic filters because Euclidean grids have directional structures (up, down, left, right).
    - Spectral ConvNets like ChebNets compute isotropic filters because there is no notion of directions on arbitrary graphs.



[1] Defferrard, Bresson, Vandergheynst, Convolutional neural networks on graphs with fast localized spectral filtering, 2016

# ChebNets for Multiple Graphs

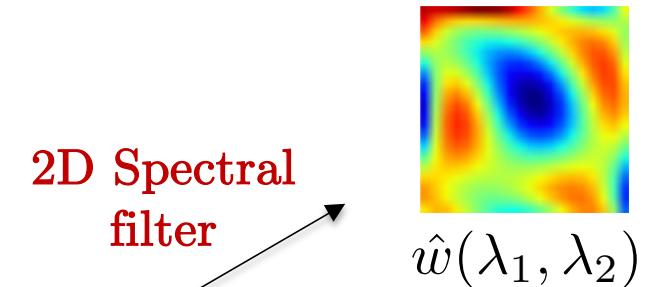
- Multi-graph spectral convolution<sup>[1]</sup> :

$$h^{\ell+1} = \eta(\hat{w}(\Delta_1, \Delta_2) * h^\ell)$$

$$h^{\ell+1} = \eta\left(\sum_{k_1=0}^{K_1-1} \sum_{k_2=0}^{K_2-1} w_{k_1, k_2} T_{k_1}(\Delta_1) T_{k_2}(\Delta_2) h^\ell\right)$$

$$\hat{w}(\lambda_1, \lambda_2) = \sum_{k_1=0}^{K_1-1} \sum_{k_2=0}^{K_2-1} w_{k_1, k_2} T_{k_1}(\Lambda_1) T_{k_2}(\Lambda_2)$$

The  $K_1, K_2$  coefficients  $w_{k_1, k_2}$  are learned by backpropagation.

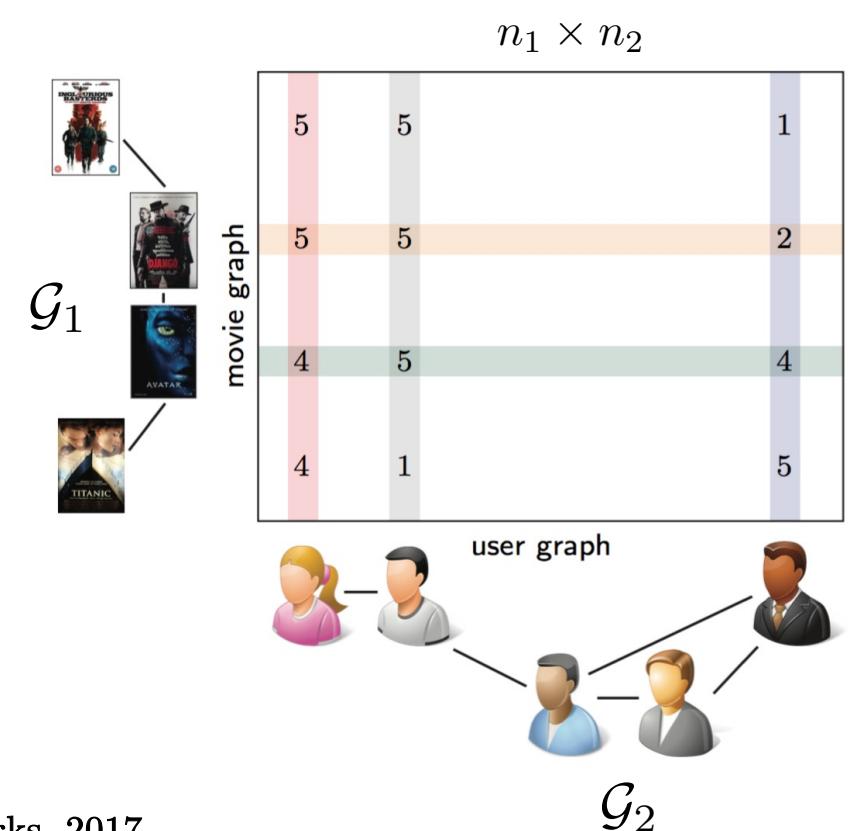


- Recommendation task (Netflix)

- Matrix Completion with Multiple Graphs<sup>[2]</sup> :

$$\min_{h \in \mathbb{R}^{n_1 \times n_2}} \|h\|_* + \mu \|\Omega \circ (h - r)\|_F^2$$

$$+ \mu_1 \underbrace{\text{tr}(h \Delta_1 h^\top)}_{\|h\|_{\mathcal{G}_1}^2} + \mu_2 \underbrace{\text{tr}(h^\top \Delta_2 h)}_{\|h\|_{\mathcal{G}_2}^2}$$



[1] F Monti, M Bronstein, X Bresson, Geometric matrix completion with recurrent multi-graph neural networks, 2017

[2] V Kalofolias, X Bresson, M Bronstein, P Vandergheynst, Matrix completion on graphs, 2014

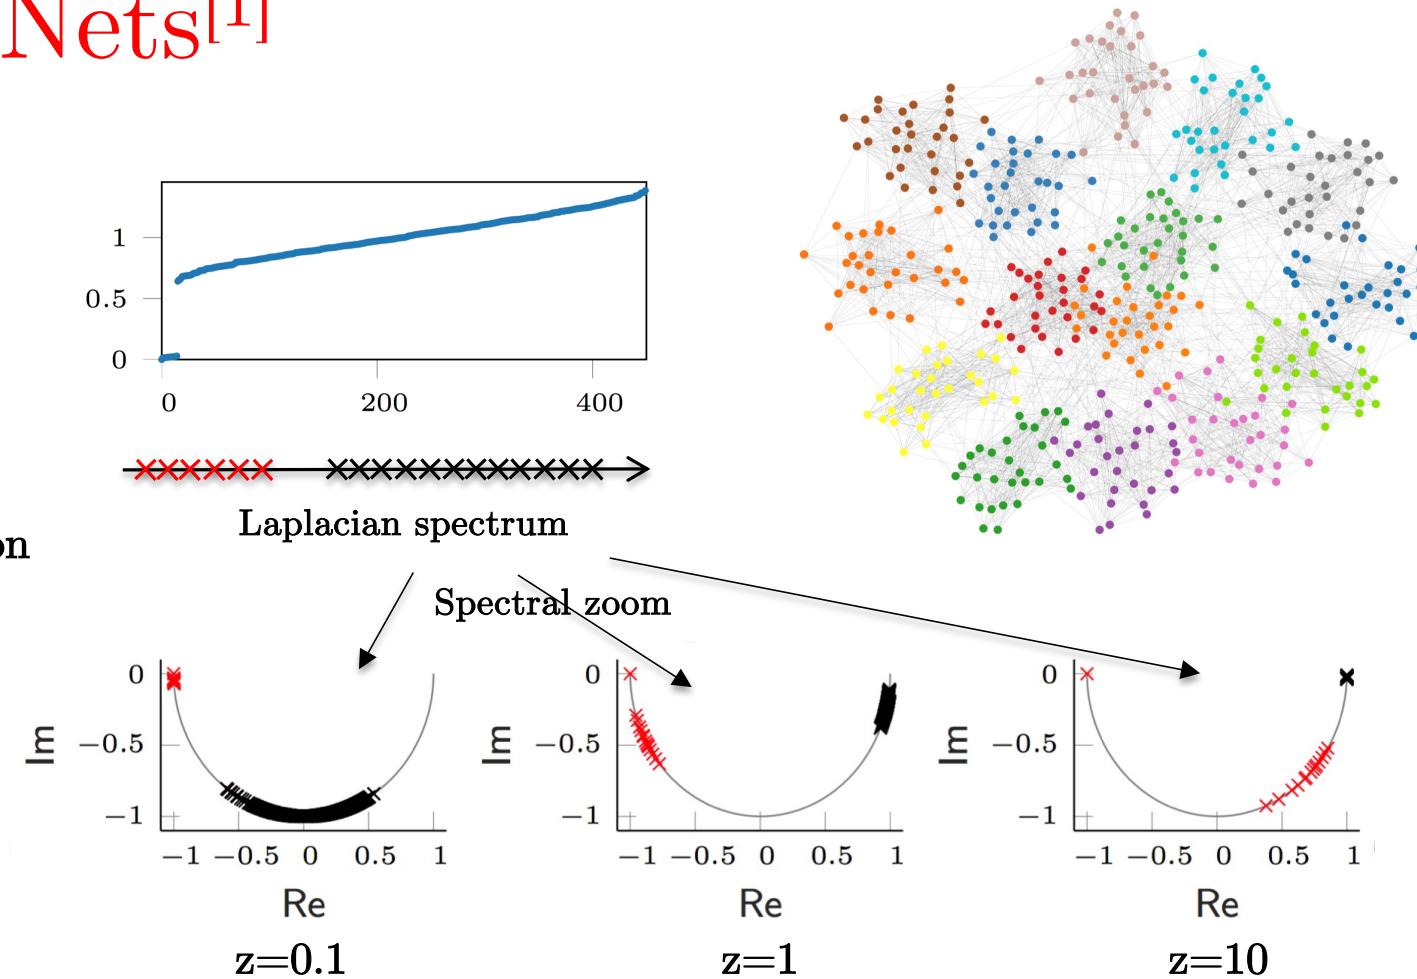
# CayleyNets[1]

- ChebGCNs are unstable to produce filters with frequency bands of interest (graph communities).
- Cayley rationals can :

Spectral zoom  $z$  focuses on frequency bands.

$$\hat{w}(\Delta) = w_0 + 2\operatorname{Re}\left\{\sum_{k=0}^{K-1} w_k \frac{(z\Delta - i)^k}{(z\Delta + i)^k}\right\}$$

The  $K$  coefficients  $w_k$  are learned by backpropagation.



- CayleyNets
  - Same properties as ChebNets
  - Localized in frequency (with spectral zoom)
  - Richer class of filters for the same order  $K$
  - Isotropic model

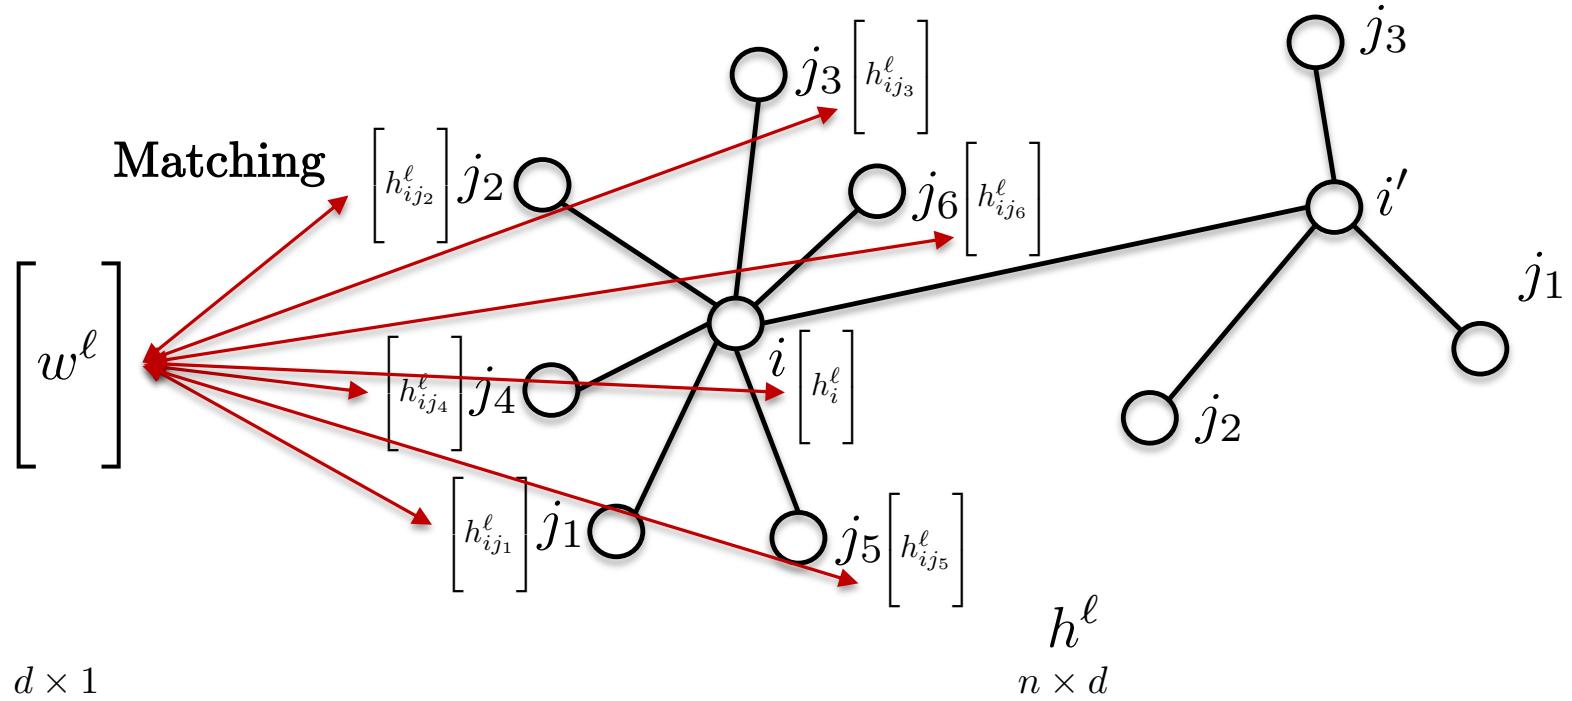
[1] R Levie, F Monti, X Bresson, MM Bronstein, CayleyNets: Graph convolutional neural networks with complex rational spectral filters, 2018

# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Template Matching

- How to define template matching for graphs ?
  - Main issue is the absence of node ordering/positioning.
  - Node indices are arbitrary and do not match the same information.
- How to design template matching invariant to node re-parametrization ?
  - Simply use the same template features for all neighbors !



$$h_i^{\ell+1} = \eta \left( \sum_{j \in \mathcal{N}_i} \underbrace{\langle w^\ell, h_{ij}^\ell \rangle}_{(h_{ij}^\ell)^T w^\ell} \right)$$

scalar  
One feature

$$h_i^{\ell+1} = \eta \left( \sum_{j \in \mathcal{N}_i} W^\ell h_{ij}^\ell \right)$$

$d$  features

$$h^{\ell+1} = \eta \left( A h^\ell W^\ell \right)$$

$n \times d$   
 $d \times d$   
Vectorial representation

# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Vanilla GCNs<sup>[1,2,3]</sup>

- Simplest formulation of spatial GCNs
  - Handle the absence of node ordering
    - Invariant by node re-parametrization
  - Deal with different neighborhood sizes
  - Local reception field by design (only neighbors are considered)
  - Weight sharing (convolution property)
  - Independent of graph size
  - Limited to isotropic capability

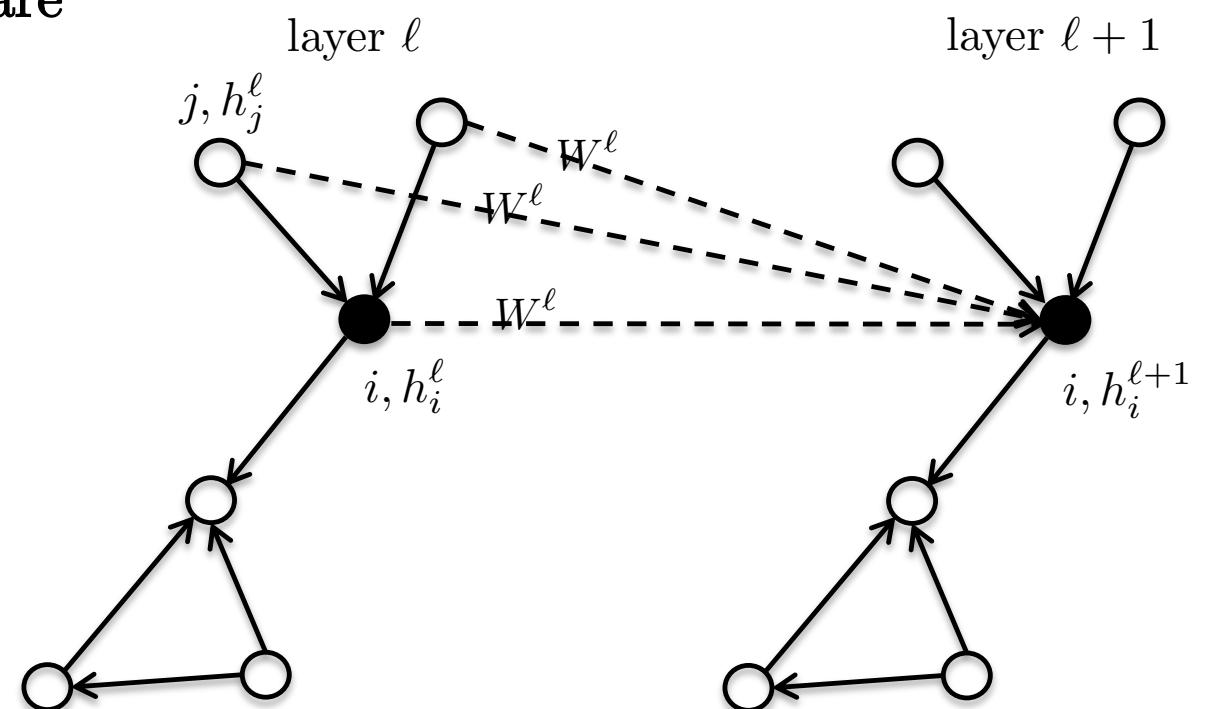
$$h^{\ell+1} = \eta(D^{-1} A h^\ell W^\ell)$$

Matrix representation

$$h_i^{\ell+1} = \eta\left(\frac{1}{d_i} \sum_{j \in \mathcal{N}_i} A_{ij} W^\ell h_j^\ell\right)$$

Vectorial representation

Mean



$$h_i^{\ell+1} = f_{\text{GCN}}(h_i^\ell, \{h_j^\ell : j \rightarrow i\})$$

[1] Scarselli, Gori, Tsoi, Hagenbuchner, Monfardini, The Graph Neural Network Model, 2009

[2] Kipf, Welling, Semi-supervised classification with graph convolutional networks, 2016

[3] Sukhbaatar, Szlam, Fergus, Learning multiagent communication with backpropagation, 2016

# ChebNets<sup>[1]</sup> and Vanilla GCNs<sup>[2,3]</sup>

- Vanilla GCNs is a simplification of ChebNets.
  - Truncated expansion of Chebyshev spectral convolution :

$$\begin{aligned} h^{\ell+1} &= \eta(w^\ell * h^\ell) \\ &= \eta(\hat{w}^\ell(\Delta)h^\ell) \\ &= \eta\left(\sum_{k=0}^{K-1} w_k^\ell T_k(\Delta)h^\ell\right) \end{aligned}$$

Suppose  $K = 2$ ,  $w_0^\ell = w^\ell$ ,  $w_1^\ell = -w^\ell$ ,  $\lambda_n = 2$ ,

$$\begin{aligned} h^{\ell+1} &= \eta(w^\ell(T_0(\Delta) - T_1(\Delta))h^\ell) \\ &= \eta(w^\ell(I + D^{-1/2}AD^{-1/2})h^\ell) \quad \text{Operator with largest eigenvalue in } [0,2] \text{ may cause divergence.} \end{aligned}$$

Add self-loop to graphs  $\hat{A} = A + I$

$$\begin{aligned} h^{\ell+1} &= \eta(w^\ell \hat{D}^{-1/2} \hat{A} \hat{D}^{-1/2} h^\ell) \quad (\text{single feature}) \\ &= \eta(\hat{D}^{-1/2} \hat{A} \hat{D}^{-1/2} h^\ell W^\ell) \quad (d \text{ features}) \end{aligned}$$

$$h_i^{\ell+1} = \eta\left(\frac{1}{\hat{d}_i^{-1/2}} \sum_{j \in \mathcal{N}_i} \frac{1}{\hat{d}_j^{-1/2}} \hat{A}_{ij} W^\ell h_j^\ell\right)$$

with

$$\begin{cases} T_0 = I \\ T_1 = \tilde{\Delta} = \frac{2}{\lambda_n} \Delta - I \stackrel{\lambda_n=2}{=} \Delta - I \\ \Delta = I - D^{-1/2}AD^{-1/2} \end{cases}$$

[1] Defferrard, Bresson, Vandergheynst, Convolutional neural networks on graphs with fast localized spectral filtering, 2016

[2] Kipf, Welling, Semi-supervised classification with graph convolutional networks, 2016

[3] Sukhbaatar, Szlam, Fergus, Learning multiagent communication with backpropagation, 2016

# GraphSage<sup>[1]</sup>

- Vanilla GCNs (supposing  $A_{ij} = 1$ ) : 
$$h_i^{\ell+1} = \eta\left(\frac{1}{d_i} \sum_{j \in \mathcal{N}_i} W^\ell h_j^\ell\right)$$
- GraphSage :
  - Differentiate template weights  $W^l$  between neighbors  $h_j$  and central node  $h_i$ .
  - Isotropic GCNs

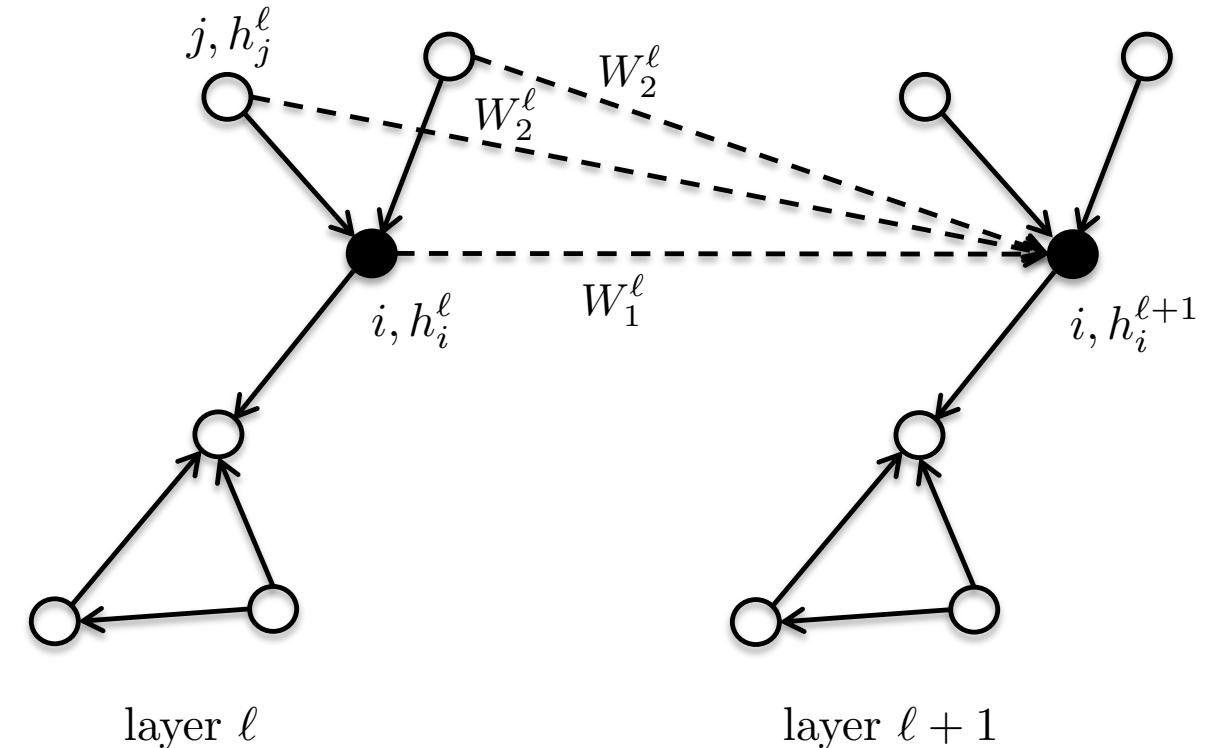
$$h_i^{\ell+1} = \eta\left( \underbrace{W_1^\ell h_i^\ell}_{d \times 1} + \frac{1}{d_i} \sum_{j \in \mathcal{N}_i} \underbrace{W_2^\ell h_j^\ell}_{d \times d} \underbrace{\sum_{j \in \mathcal{N}_i} h_j^\ell}_{d \times 1} \right)$$

Mean <sub>$j \in \mathcal{N}_i$</sub>   $W_2^\ell h_j^\ell$

Or alternatively,

$$\text{Max}_{j \in \mathcal{N}_i} W_2^\ell h_j^\ell$$

$$\text{LSTM}^\ell(h_j^\ell)$$



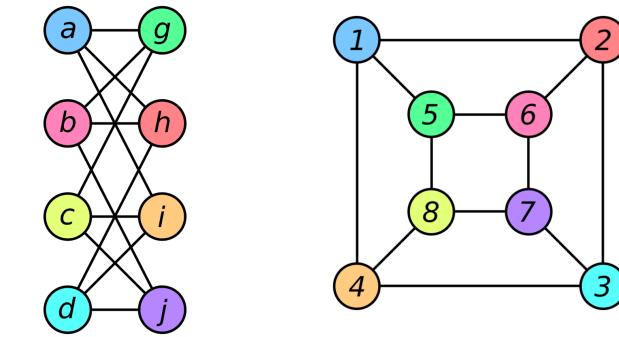
[1] Hamilton, Ying, Leskovec, Inductive representation learning on large graphs, 2017

# Graph Isomorphism Networks<sup>[1]</sup> (GIN)

- Architecture that can differentiate graphs that are not isomorphic.
  - Graph isomorphism is an equivalent relation for similar graph structures.
- Isotropic GCNs

$$h_i^{\ell+1} = \text{ReLU}(W_2^\ell \text{ReLU}(\text{BN}(\hat{h}_i^{\ell+1})) )$$
$$\hat{h}_i^{\ell+1} = (1 + \epsilon) h_i^\ell + \sum_{j \in \mathcal{N}_i} h_j^\ell$$

Batch  
normalization



Example of two  
isomorphic graphs

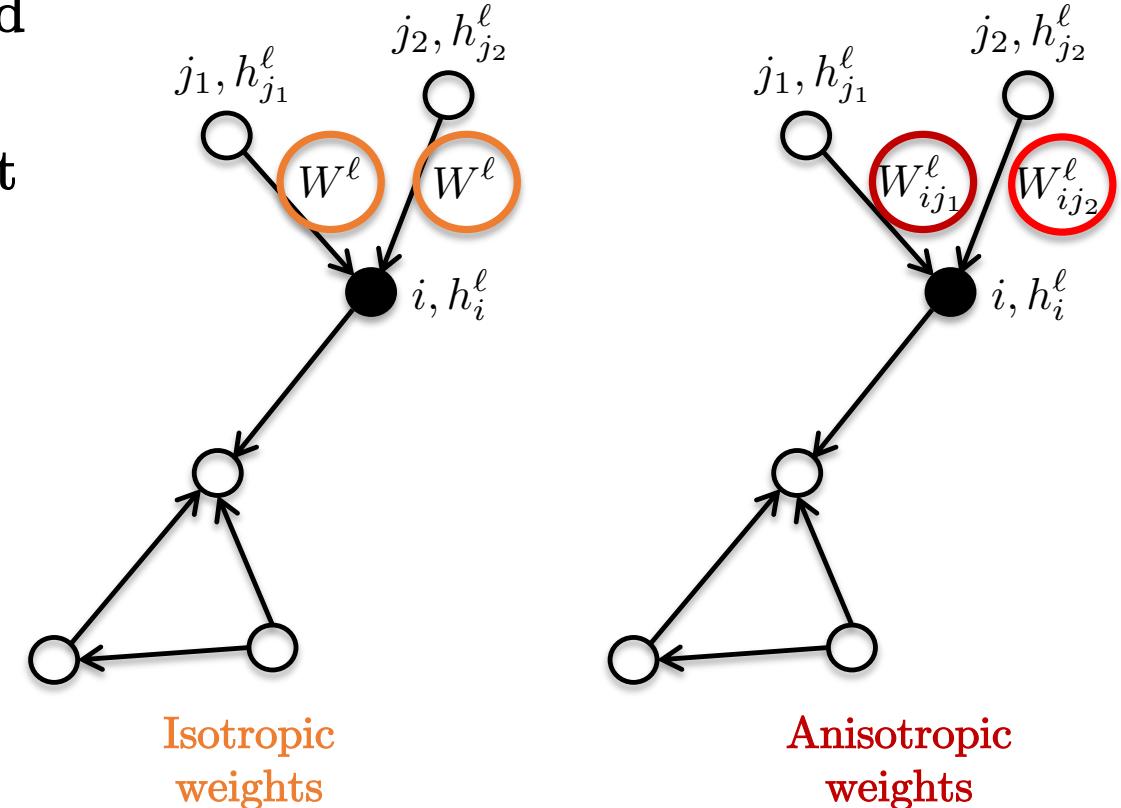
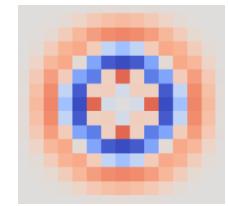
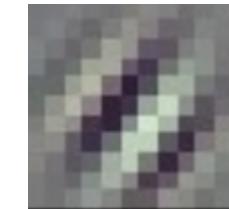
[1] Xu, Hu, Leskovec, Jegelka, How powerful are graph neural networks?, 2018

# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - **Anisotropic GCNs**
  - DGL : Implemented layers and lab demo
- Conclusion

# Anisotropic GCNs

- Reminder :
  - Standard ConvNets produce anisotropic filters because Euclidean grids have directional structures (up, down, left, right).
  - GCNs such as ChebNets, CayleyNets, Vanilla GCNs, GraphSage, GIN compute isotropic filters as there is no notion of directions on arbitrary graphs.
- How to get anisotropy back in GNNs ?
  - Natural edge features<sup>[1,2]</sup> if available (e.g. different bond connections between atoms).
  - We need an anisotropic mechanism that is independent of the node parametrization.
  - Edge degrees<sup>[3]</sup>/Edge gates<sup>[4]</sup>/Attention mechanism<sup>[5]</sup> : MoNets<sup>[3]</sup>, GAT<sup>[5]</sup>, GatedGCNs<sup>[4]</sup> can treat neighbors differently.



[1] Gilmer, Schoenholz, Riley, Vinyals, Dahl, Neural message passing for quantum chemistry, 2017

[2] Bresson, Laurent, A Two-Step Graph Convolutional Decoder for Molecule Generation, 2019

[3] Monti, Boscaini, Masci, Rodolà, J. Svoboda, M. Bronstein, Geometric deep learning on graphs and manifolds using mixture model CNNs, 2016

[4] Bresson, Laurent, Residual gated graph convnets, 2017

[5] Velickovic, Cucurull, Casanova, Romero, Lio, Bengio, Graph Attention Networks, 2018

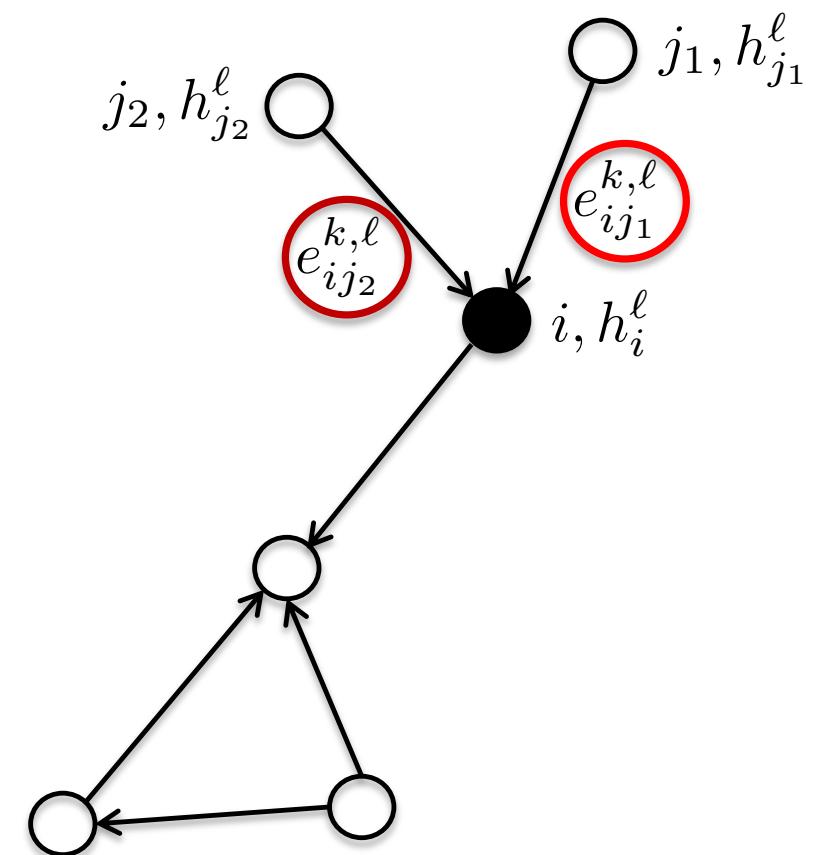
# MoNets<sup>[1]</sup>

- MoNets<sup>[1]</sup> leverage the Bayesian Gaussian Mixture Model (GMM)<sup>[2]</sup>.

$$h_i^{\ell+1} = \text{ReLU} \left( \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} e_{ij}^{k,\ell} W_1^{k,\ell} h_j^\ell \right)$$

$$e_{ij}^{k,\ell} = \exp \left( -\frac{1}{2} (u_{ij}^\ell - \mu_k)^\top (\Sigma_k)_{2 \times 2}^{-1} (u_{ij}^\ell - \mu_k) \right)$$

$$u_{ij}^\ell = \text{Tanh} \left( A^\ell \begin{pmatrix} \deg_i^{-1/2} \\ \deg_j^{-1/2} \end{pmatrix} + a^\ell \right)$$



[1] Monti, Boscaini, Masci, Rodolà, J. Svoboda, M. Bronstein, Geometric deep learning on graphs and manifolds using mixture model CNNs, 2016

[2] Dempster, Laird, Rubin, Maximum Likelihood from Incomplete Data Via the EM Algorithm, 1977

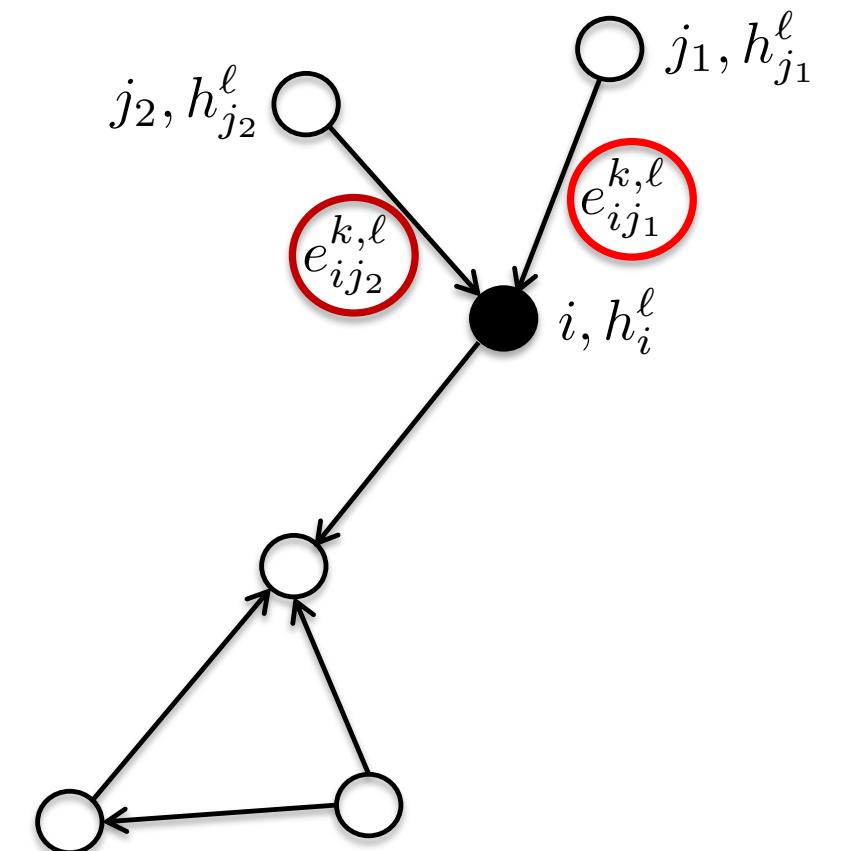
# Graph Attention Networks<sup>[1]</sup> (GAT)

- GAT uses the attention mechanism<sup>[2]</sup> to introduce anisotropy in the neighborhood aggregation function.
- The network employs a multi-headed architecture to increase the learning capacity, similar to Transformers<sup>[3]</sup>.

$$h_i^{\ell+1} = \text{Concat}_{k=1}^K \left( \text{ELU} \left( \sum_{j \in \mathcal{N}_i} e_{ij}^{k,\ell} W_1^{k,\ell} h_j^\ell \right) \right)$$

$$e_{ij}^{k,\ell} = \text{Softmax}_{\mathcal{N}_i}(\hat{e}_{ij}^{k,\ell}) = \frac{\exp(\hat{e}_{ij}^{k,\ell})}{\sum_{j' \in \mathcal{N}_i} \exp(\hat{e}_{ij'}^{k,\ell})}$$

$$\hat{e}_{ij}^{k,\ell} = \text{LeakyReLU} \left( W_2^{k,\ell} \underbrace{\text{Concat} \left( W_1^{k,\ell} h_i^\ell, W_1^{k,\ell} h_j^\ell \right)}_{\frac{2d}{K} \times 1} \right)$$



[1] Velickovic, Cucurull, Casanova, Romero, Lio, Bengio, Graph Attention Networks, 2018

[2] Bahdanau, Cho, Bengio, Neural machine translation by jointly learning to align and translate, 2014

[3] Vaswani, Shazeer, Parmar, Uszkoreit, Jones, Gomez, Kaiser, Polosukhin, Attention is all you need, 2017

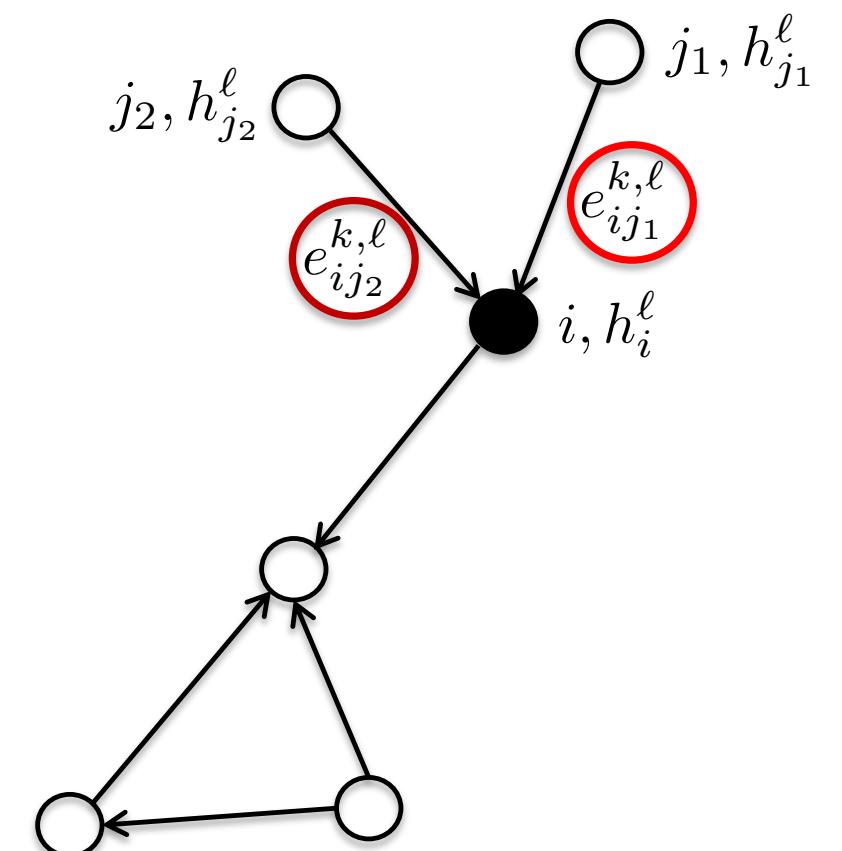
# Gated Graph ConvNets<sup>[1]</sup>

- GatedGCNs use edge gates to design an anisotropic variant of GCNs.
  - Edge gates can be regarded as a soft attention process, related to the standard sparse attention mechanism<sup>[2]</sup>.
  - Edge features are explicit (important for edge prediction tasks).
- Residual connections and batch normalization enhance learning speed and generalization.

$$h_i^{\ell+1} = h_i^\ell + \text{ReLU}\left(\text{BN}\left(W_1^\ell h_i^\ell + \sum_{j \in \mathcal{N}_i} e_{ij}^\ell \odot W_2^\ell h_j^\ell\right)\right)$$

$$e_{ij}^\ell = \frac{\sigma(\hat{e}_{ij}^\ell)}{\sum_{j' \in \mathcal{N}_i} \sigma(\hat{e}_{ij'}^\ell) + \varepsilon}$$

$$\hat{e}_{ij}^\ell = \hat{e}_{ij}^{\ell-1} + \text{ReLU}\left(\text{BN}\left(V_1^\ell h_i^{\ell-1} + V_2^\ell h_j^{\ell-1} + V_3^\ell \hat{e}_{ij}^{\ell-1}\right)\right)$$



[1] Bresson, Laurent, Residual gated graph convnets, 2017

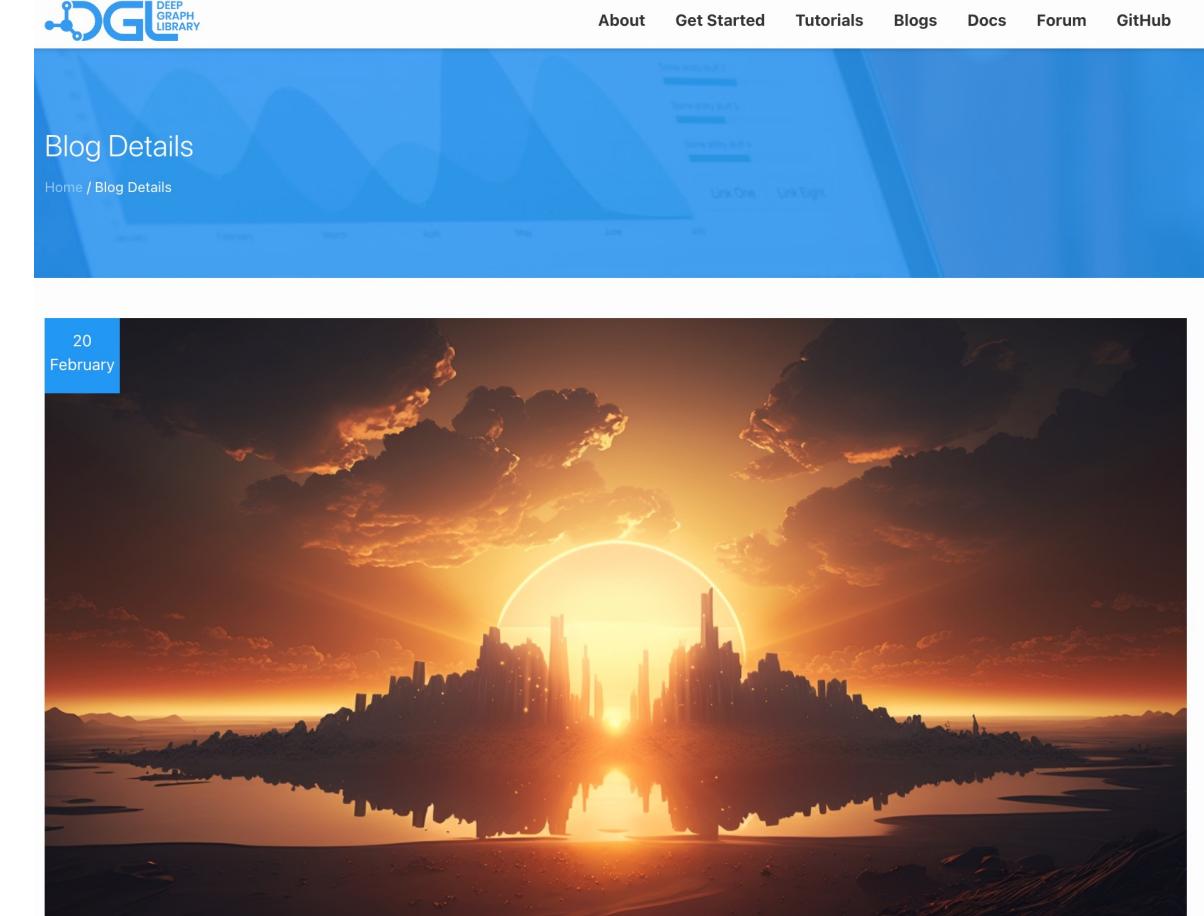
[2] Bahdanau, Cho, Bengio, Neural machine translation by jointly learning to align and translate, 2014

# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

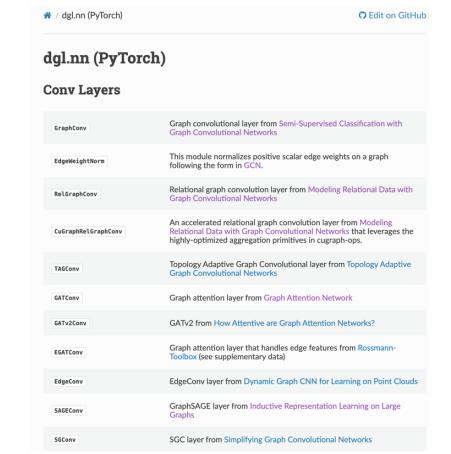
# DGL Library (Amazon)

- Website : <https://www.dgl.ai>, <https://docs.dgl.ai>
- DGL 1.0 : <https://www.dgl.ai/release/2023/02/20/release.html>



# DGL-implemented GCNs

- Implemented layers in DGL : <https://docs.dgl.ai/api/python/nn-pytorch.html>
- Graph convolutional layers
  - ChebNet : <https://docs.dgl.ai/generated/dgl.nn.pytorch.conv.ChebConv.html?highlight=chebnet>
  - GCN :  
<https://docs.dgl.ai/generated/dgl.nn.pytorch.conv.GraphConv.html#dgl.nn.pytorch.conv.GraphConv>
  - GraphSage :  
<https://docs.dgl.ai/generated/dgl.nn.pytorch.conv.SAGEConv.html#dgl.nn.pytorch.conv.SAGEConv>
  - GIN : <https://docs.dgl.ai/generated/dgl.nn.pytorch.conv.GINConv.html#dgl.nn.pytorch.conv.GINConv>
- Attention layers
  - GAT :  
<https://docs.dgl.ai/generated/dgl.nn.pytorch.conv.GATConv.html#dgl.nn.pytorch.conv.GATConv>
  - GraphTransformer : [https://doc-build.dgl.ai/notebooks/sparse/graph\\_transformer.html](https://doc-build.dgl.ai/notebooks/sparse/graph_transformer.html)



# Lab 1: GatedGCNs<sup>[1]</sup>

- Run code01.ipynb

jupyter code01 Last Checkpoint: a minute ago (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel) Logout

### Graph Convolutional Networks

#### GatedGCNs

Residual Gated Graph ConvNets, X Bresson, T Laurent, ICLR 2018, [arXiv:1711.07553](https://arxiv.org/abs/1711.07553)

```
In [1]: # For Google Colaboratory
import sys
os
if 'google.colab' in sys.modules:
    # mount google drive
    from google.colab import drive
    drive.mount('/content/gdrive')
    path_to_file = '/content/gdrive/My Drive/CS6208_codes/codes/labs_lecture05'
    print(path_to_file)
    # change current path to the folder containing "path_to_file"
    os.chdir(path_to_file)
!pwd
```

### Train GNN

```
In [9]: # datasets
train_loader = DataLoader(trainset, batch_size=50, shuffle=True, collate_fn=collate)
test_loader = DataLoader(testset, batch_size=50, shuffle=False, collate_fn=collate)
val_loader = DataLoader(valset, batch_size=50, shuffle=False, drop_last=False, collate_fn=collate)

# Create model
net_parameters = {}
net_parameters['input_dim'] = 1
net_parameters['hidden_dim'] = 100
net_parameters['output_dim'] = 8 # nb of classes
net_parameters['L'] = 4
net = GatedGCN_Net(net_parameters)

optimizer = torch.optim.Adam(net.parameters(), lr=0.0001)

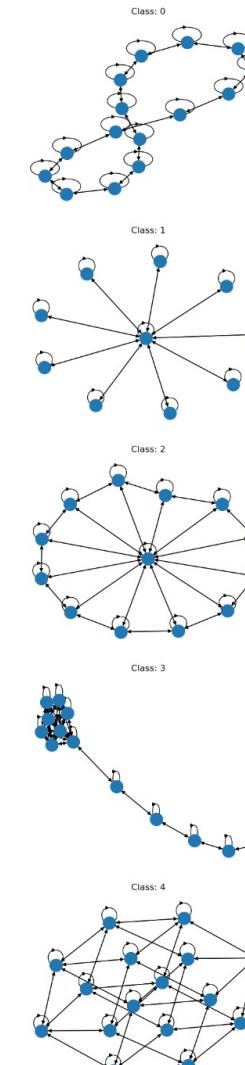
epoch_train_losses = []
epoch_val_losses = []
epoch_train_accs = []
epoch_val_accs = []
for epoch in range(50):

    start = time.time()
    epoch_train_loss, epoch_train_acc = train_one_epoch(net, train_loader)
    epoch_val_loss, epoch_val_acc = evaluate_network(net, val_loader)
    epoch_val_loss, epoch_val_acc = evaluate_network(net, val_loader)

    print('Epoch {}, time {:.4f}, train_loss: {:.4f}, test_loss: {:.4f}, val_loss: {:.4f} \n'.format(epoch, time.time() - start, epoch_train_loss, epoch_val_loss, epoch_val_loss))

Epoch 0, time 1.6418, train_loss: 2.0092, test_loss: 2.0319, val_loss: 2.0319
Epoch 1, time 1.6794, train_loss: 1.9021, test_loss: 1.9863, val_loss: 1.9863
Epoch 2, time 1.7416, train_loss: 1.8032, test_loss: 1.9217, val_loss: 1.9217
Epoch 3, time 1.7690, train_loss: 1.7355, test_loss: 1.8438, val_loss: 1.8438
Epoch 4, time 1.7733, train_loss: 1.6713, test_loss: 1.7568, val_loss: 1.7568
Epoch 5, time 1.4200, train_loss: 0.5167, test_loss: 0.4712, val_loss: 0.4712
Epoch 6, time 1.7266, train_loss: 1.5380, test_loss: 1.5942, val_loss: 1.5942
Epoch 7, time 1.5330, train_loss: 1.4817, test_loss: 1.5212, val_loss: 1.5212
Epoch 8, time 1.4076, train_loss: 1.4333, test_loss: 1.4440, val_loss: 1.4440
Epoch 9, time 1.6103, train_loss: 1.3739, test_loss: 1.3811, val_loss: 1.3811
    train_acc: 0.5020, test_acc: 0.3700, val_acc: 0.3700
    train_acc: 0.5620, test_acc: 0.4300, val_acc: 0.4300
    train_acc: 0.5920, test_acc: 0.4600, val_acc: 0.4600
    train_acc: 0.6220, test_acc: 0.4900, val_acc: 0.4900
    train_acc: 0.6520, test_acc: 0.5200, val_acc: 0.5200
    train_acc: 0.6820, test_acc: 0.5500, val_acc: 0.5500
    train_acc: 0.7120, test_acc: 0.5800, val_acc: 0.5800
    train_acc: 0.7420, test_acc: 0.6100, val_acc: 0.6100
    train_acc: 0.7720, test_acc: 0.6400, val_acc: 0.6400
    train_acc: 0.8020, test_acc: 0.6700, val_acc: 0.6700
    train_acc: 0.8320, test_acc: 0.7000, val_acc: 0.7000
    train_acc: 0.8620, test_acc: 0.7300, val_acc: 0.7300
    train_acc: 0.8920, test_acc: 0.7600, val_acc: 0.7600
    train_acc: 0.9220, test_acc: 0.7900, val_acc: 0.7900
    train_acc: 0.9520, test_acc: 0.8200, val_acc: 0.8200
    train_acc: 0.9820, test_acc: 0.8500, val_acc: 0.8500
    train_acc: 1.0000, test_acc: 0.8800, val_acc: 0.8800
```

[1] Bresson, Laurent, Residual gated graph convnets, 2017



# Lab 1: GatedGCNs

DGL creates a batch of graphs

Create artificial node feature (input degree) and edge feature (value 1)

```
# collate function
def collate(samples):
    graphs, labels = map(list, zip(*samples)) # samples is a list of pairs (graph, label).
    labels = torch.tensor(labels)
    tab_sizes_n = [ graph.number_of_nodes() for i in range(len(graphs)) ] # graph sizes
    tab_snorm_n = [ torch.FloatTensor(size,1).fill_(1./float(size)) for size in tab_sizes_n ]
    snorm_n = torch.cat(tab_snorm_n).sqrt() # normalization constant for better optimization
    tab_sizes_e = [ graph.number_of_edges() for i in range(len(graphs)) ] # nb of edges
    tab_snorm_e = [ torch.FloatTensor(size,1).fill_(1./float(size)) for size in tab_sizes_e ]
    snorm_e = torch.cat(tab_snorm_e).sqrt() # normalization constant for better optimization
    batched_graph = dgl.batch(graphs) # batch graphs
    return batched_graph, labels, snorm_n, snorm_e

# create artificial data feature (= in degree) for each node
def create_artificial_features(dataset):
    for (graph,_) in dataset:
        graph.ndata['feat'] = graph.in_degrees().view(-1, 1).float()
        graph.edata['feat'] = torch.ones(graph.number_of_edges(),1)
    return dataset

# use artificial graph dataset of DGL
trainset = MiniGCDataset(8, 10, 20)
trainset = create_artificial_features(trainset)
print(trainset[0])

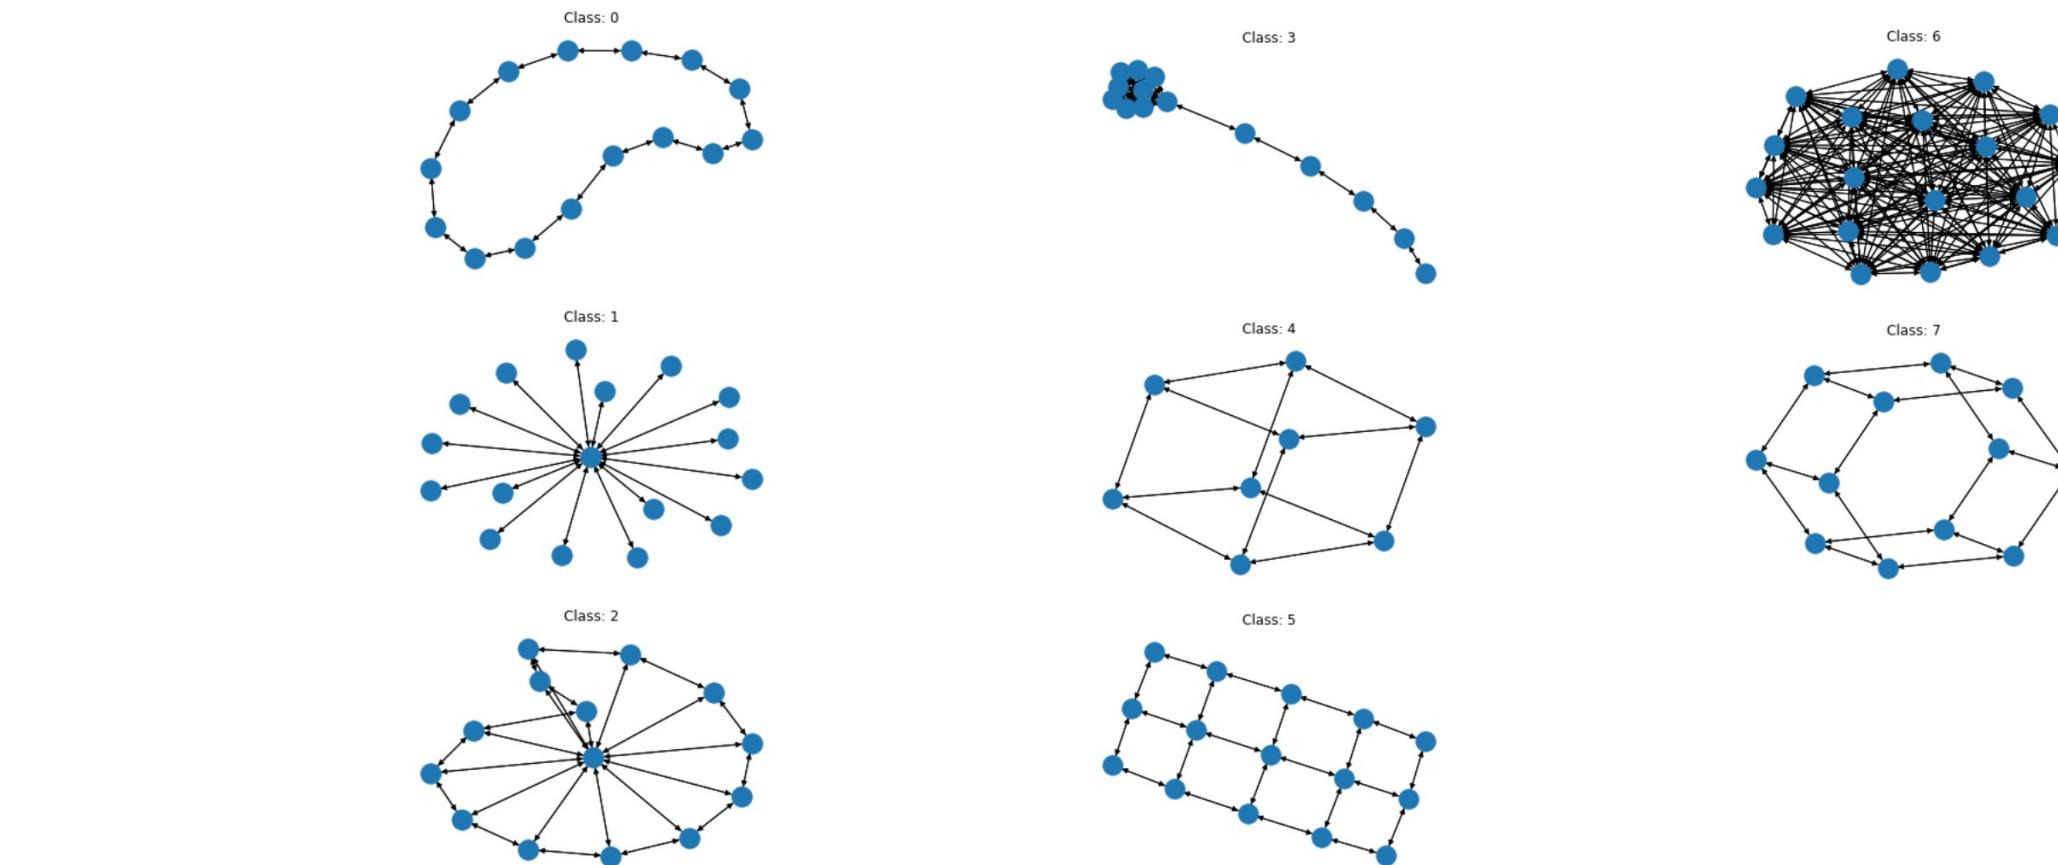
(DGLGraph(num_nodes=13, num_edges=39,
          ndata_schemes={'feat': Scheme(shape=(1,), dtype=torch.float32)}
          edata_schemes={'feat': Scheme(shape=(1,), dtype=torch.float32)}), 0)
```

$$\frac{1}{\sqrt{V_k}}, \frac{1}{\sqrt{E_k}}$$

Graph normalization constants

# Lab 1: GatedGCNs

```
# use artificial graph dataset of DGL  
trainset = MiniGCNDataset(8, 10, 20)
```

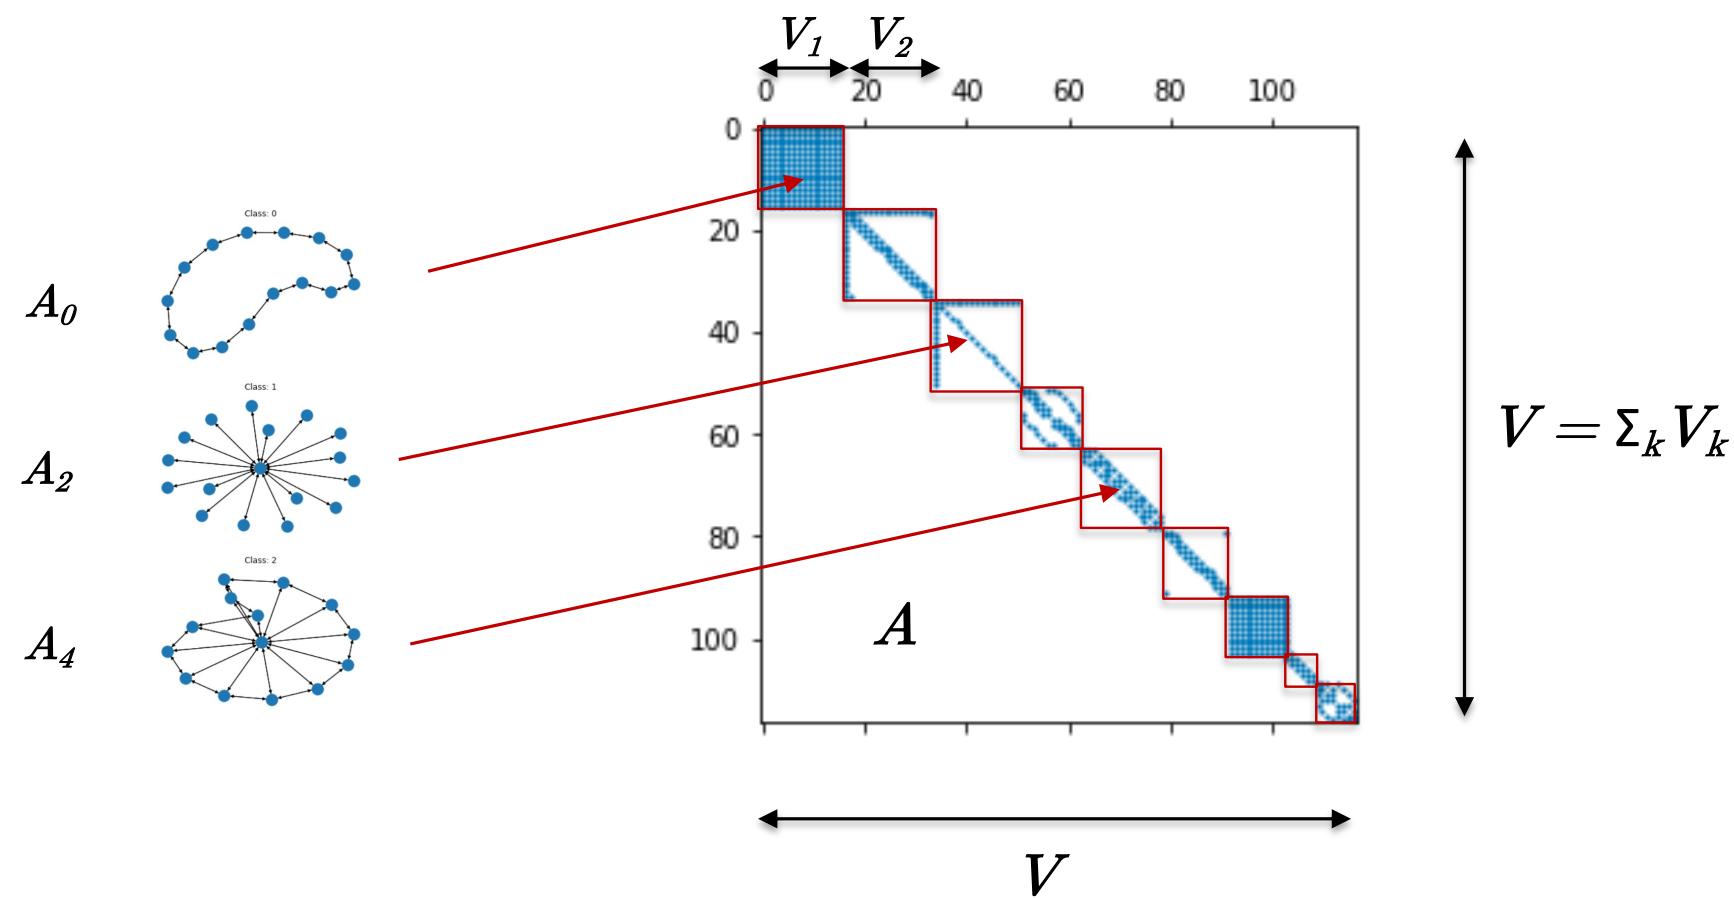


# Lab 1: GatedGCNs

- Understanding DGL :

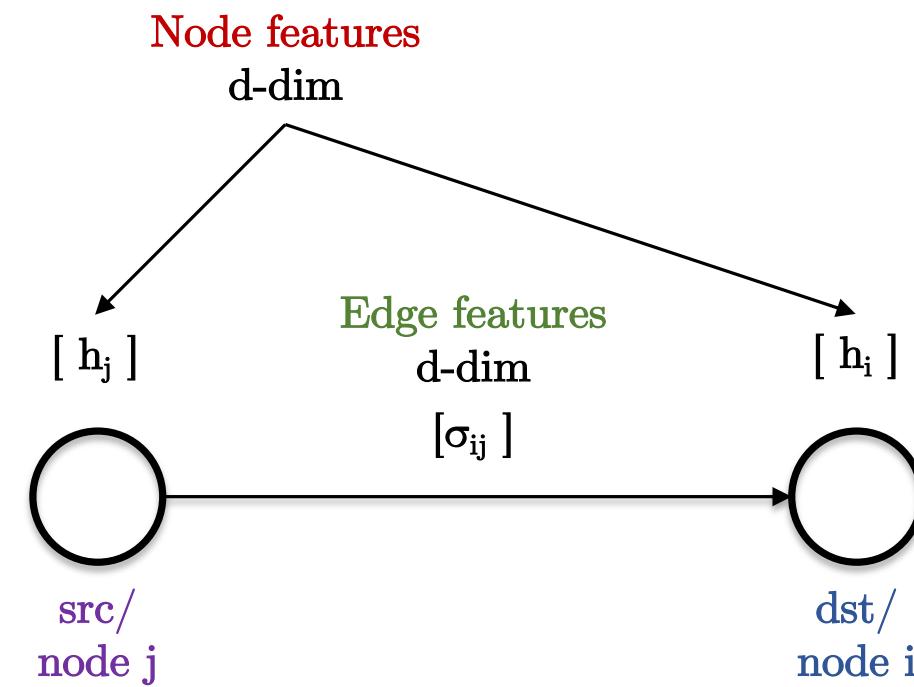
- How to process  $K$  graphs of different sizes ?

Form a (big) sparse block diagonal matrix  $A$  with  $K$  adjacency matrices  $A_k$ .



# Lab 1: GatedGCNs

- Basic structure of an edge in DGL :

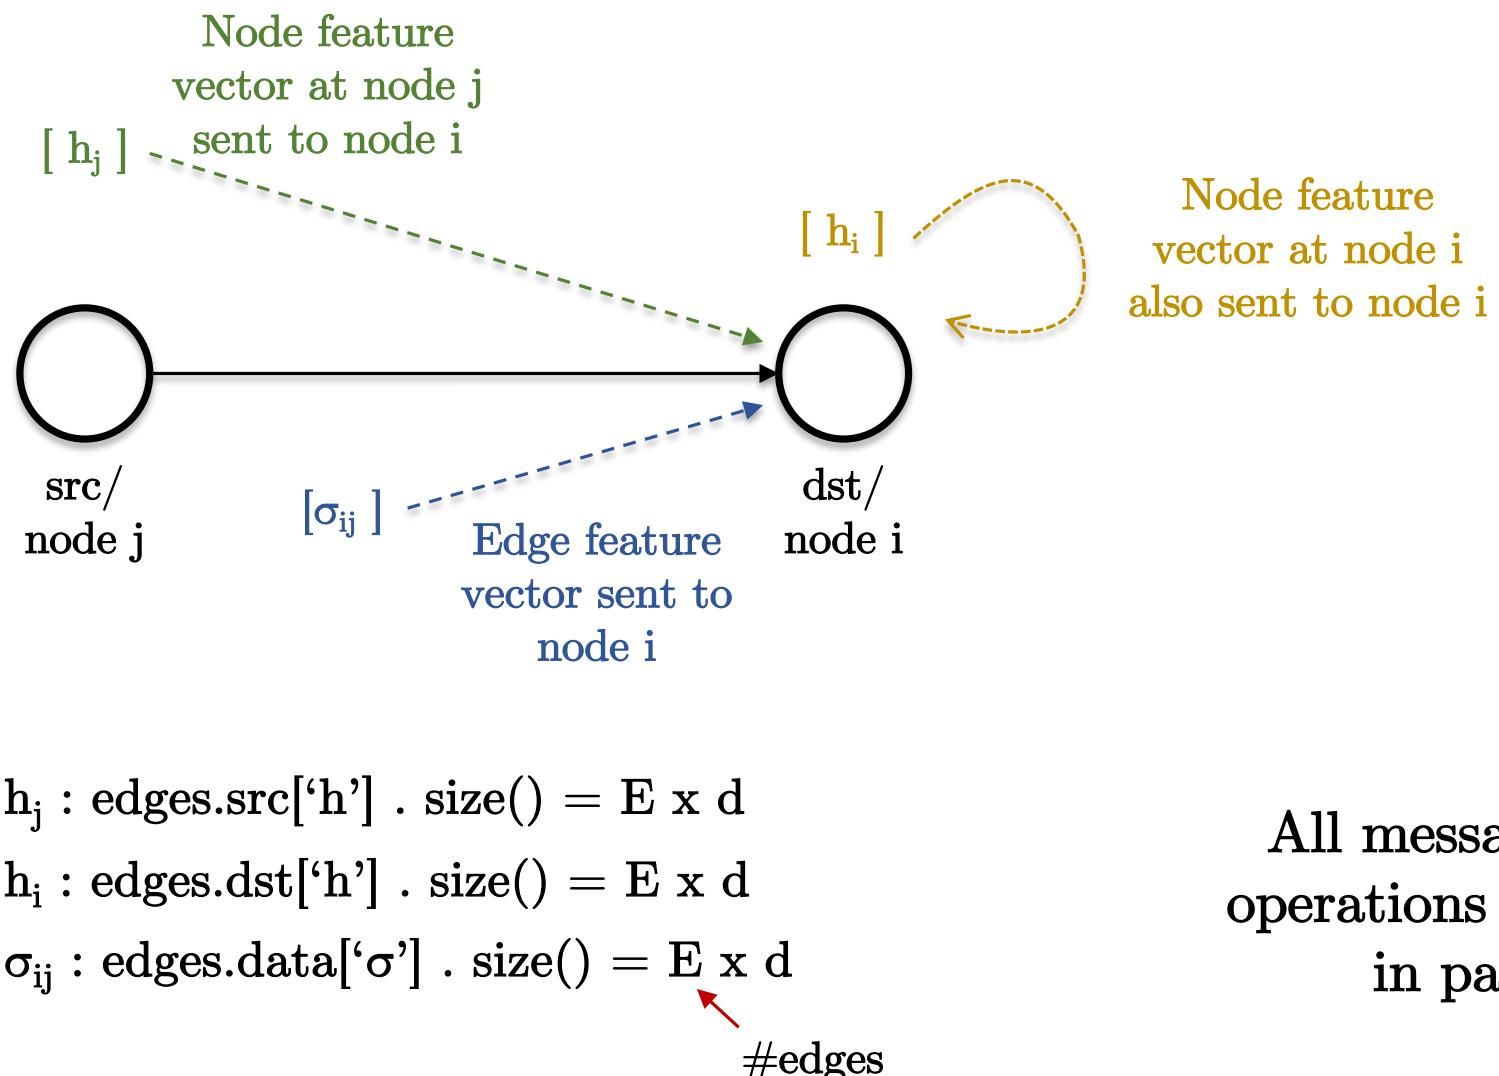


- Goal is to compute efficiently expressions of the form :

$$f_i = h_i + \sum_{j \rightarrow i} \sigma_{ij} \circ h_j$$

# Lab 1: GatedGCNs

- Step 1 : Message passing function defined on edges
  - Node feature and edge feature are passed along all edges connecting a node.



# Lab 1: GatedGCNs

- Step 2 : Reduce function defined on nodes

- Reduce functions of the form :  $f_i = h_i + \sum_{j \rightarrow i} \sigma_{ij} \circ h_j$
- Reduce function collects all messages passed in Step 1.  $\xrightarrow{\text{#nodes}}$
- Code :  $f = h_i + \text{torch.sum}(h_j \times \sigma_{ij}, \text{dim}=1) . \text{size()} = V \times d$   
 $\xrightarrow{\text{Sum over neighbors}}$

- GPU acceleration :

- DGL batches the nodes with the same number of neighbors.

$$h_j = \text{nodes.mailbox}[h_j] = \left[ \begin{array}{c} \text{batch}_1 . \text{size}() = 11 \times 12 \times d \\ \vdots \\ \text{batch}_{34} . \text{size}() = 14 \times 9 \times d \end{array} \right]$$

#nodes in batch<sub>1</sub>  
batch<sub>1</sub>.size() = 11 x 12 x d  
⋮  
batch<sub>34</sub>.size() = 14 x 9 x d  
#neighbors

$$\sigma_{ij} = \text{nodes.mailbox}[\sigma_{ij}] = \quad \text{same structure than } h_j$$

$$h_i = \text{nodes.data}[h] . \text{size()} = V \times d$$

# Lab 1: GatedGCNs

GatedGCN layer :

$$h_i^{\ell+1} = h_i^\ell + \text{ReLU}\left(\text{BN}\left(A^\ell h_i^\ell + \sum_{j \sim i} \eta(e_{ij}^\ell) \odot B^\ell h_j^\ell\right)\right)/\sqrt{V_k},$$

$$\eta(e_{ij}^\ell) = \frac{\sigma(e_{ij}^\ell)}{\sum_{j' \sim i} \sigma(e_{ij'}^\ell) + \varepsilon},$$

$$e_{ij}^{\ell+1} = e_{ij}^\ell + \text{ReLU}\left(\text{BN}\left(C^\ell e_{ij}^\ell + D^\ell h_i^{\ell+1} + E^\ell h_j^{\ell+1}\right)\right)/\sqrt{E_k}.$$

```

class GatedGCN_layer(nn.Module):

    def __init__(self, input_dim, output_dim):
        super(GatedGCN_layer, self).__init__()
        self.A = nn.Linear(input_dim, output_dim, bias=True)
        self.B = nn.Linear(input_dim, output_dim, bias=True)
        self.C = nn.Linear(input_dim, output_dim, bias=True)
        self.D = nn.Linear(input_dim, output_dim, bias=True)
        self.E = nn.Linear(input_dim, output_dim, bias=True)
        self.bn_node_h = nn.BatchNorm1d(output_dim)
        self.bn_node_e = nn.BatchNorm1d(output_dim)

    def message_func(self, edges):
        Bh_j = edges.src['Bh']
        e_ij = edges.data['Ce'] + edges.src['Dh'] + edges.dst['Eh'] # e_ij = Ce_ij + Dhi + Ehj
        edges.data['e'] = e_ij
        return {'Bh_j': Bh_j, 'e_ij': e_ij}

    def reduce_func(self, nodes):
        Ah_i = nodes.data['Ah']
        Bh_j = nodes.mailbox['Bh_j']
        e = nodes.mailbox['e_ij']
        sigma_ij = torch.sigmoid(e) # sigma_ij = sigmoid(e_ij)
        h = Ah_i + torch.sum(sigma_ij * Bh_j, dim=1) / torch.sum(sigma_ij, dim=1) # hi = Ahi + sum_j eta_ij * Bj
        return {'h': h}

    def forward(self, g, h, e, snorm_n, snorm_e):

        h_in = h # residual connection
        e_in = e # residual connection

        g.ndata['h'] = h
        g.ndata['Ah'] = self.A(h)
        g.ndata['Bh'] = self.B(h)
        g.ndata['Dh'] = self.D(h)
        g.ndata['Eh'] = self.E(h)
        g.edata['e'] = e
        g.edata['Ce'] = self.C(e)
        g.update_all(self.message_func, self.reduce_func)
        h = g.ndata['h'] # result of graph convolution
        e = g.edata['e'] # result of graph convolution

        h = h * snorm_n # normalize activation w.r.t. graph node size
        e = e * snorm_e # normalize activation w.r.t. graph edge size

        h = self.bn_node_h(h) # batch normalization
        e = self.bn_node_e(e) # batch normalization

        h = F.relu(h) # non-linear activation
        e = F.relu(e) # non-linear activation

        h = h_in + h # residual connection
        e = e_in + e # residual connection

        return h, e

```

# Lab 1: GatedGCNs

MLP classifier layer

```
class MLP_layer(nn.Module):

    def __init__(self, input_dim, output_dim, L=2): # L = nb of hidden layers
        super(MLP_layer, self).__init__()
        list_FC_layers = [ nn.Linear( input_dim, input_dim, bias=True ) for l in range(L) ]
        list_FC_layers.append(nn.Linear( input_dim, output_dim , bias=True ))
        self.FC_layers = nn.ModuleList(list_FC_layers)
        self.L = L

    def forward(self, x):
        y = x
        for l in range(self.L):
            y = self.FC_layers[l](y)
            y = F.relu(y)
        y = self.FC_layers[self.L](y)
        return y
```

GatedGCN Network

Node input embedding

Edge input embedding

Run graphNN layers

Compute graph vectorial representation by a (simple) average of all node features with DGL.

Use MLP classifier

```
class GatedGCN_Net(nn.Module):

    def __init__(self, net_parameters):
        super(GatedGCN_Net, self).__init__()
        input_dim = net_parameters['input_dim']
        hidden_dim = net_parameters['hidden_dim']
        output_dim = net_parameters['output_dim']
        L = net_parameters['L']
        self.embedding_h = nn.Linear(input_dim, hidden_dim)
        self.embedding_e = nn.Linear(1, hidden_dim)
        self.GatedGCN_layers = nn.ModuleList([ GatedGCN_layer(hidden_dim, hidden_dim) for _ in range(L) ])
        self.MLP_layer = MLP_layer(hidden_dim, output_dim)

    def forward(self, g, h, e, snorm_n, snorm_e):

        # input embedding
        h = self.embedding_h(h)
        e = self.embedding_e(e)

        # graph convnet layers
        for GGCN_layer in self.GatedGCN_layers:
            h,e = GGCN_layer(g,h,e,snorm_n,snorm_e)

        # MLP classifier
        g.ndata['h'] = h
        y = dgl.mean_nodes(g, 'h')
        y = self.MLP_layer(y)

        return y
```

# Lab 1: GatedGCNs

```
def train_one_epoch(net, data_loader):
    """
    train one epoch
    """
    net.train()
    epoch_loss = 0
    epoch_train_acc = 0
    nb_data = 0
    gpu_mem = 0
    for iter, (batch_graphs, batch_labels, batch_snorm_n, batch_snorm_e) in enumerate(data_loader):
        batch_x = batch_graphs.ndata['feat'].to(device)
        batch_e = batch_graphs.edata['feat'].to(device)
        batch_snorm_n = batch_snorm_n.to(device)
        batch_snorm_e = batch_snorm_e.to(device)
        batch_labels = batch_labels.to(device)
        batch_scores = net.forward(batch_graphs, batch_x, batch_e, batch_snorm_n, batch_snorm_e)
        gpu_mem = net.gpu_memory(gpu_mem)
        loss = net.loss(batch_scores, batch_labels)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        epoch_loss += loss.detach().item()
        epoch_train_acc += net.accuracy(batch_scores, batch_labels)
        nb_data += batch_labels.size(0)
    epoch_loss /= (iter + 1)
    epoch_train_acc /= nb_data
    return epoch_loss, epoch_train_acc, gpu_mem
```

## Train function

```
# datasets
train_loader = DataLoader(trainset, batch_size=50, shuffle=True, collate_fn=collate)
test_loader = DataLoader(testset, batch_size=50, shuffle=False, collate_fn=collate)
val_loader = DataLoader(valset, batch_size=50, shuffle=False, drop_last=False, collate_fn=collate)

# Create model
net_parameters = {}
net_parameters['input_dim'] = 1
net_parameters['hidden_dim'] = 100
net_parameters['output_dim'] = 8 # nb of classes
net_parameters['L'] = 4
net = GatedGCN_Net(net_parameters)
net = net.to(device)

optimizer = torch.optim.Adam(net.parameters(), lr=0.0005)

epoch_train_losses = []
epoch_test_losses = []
epoch_val_losses = []
epoch_train_accs = []
epoch_test_accs = []
epoch_val_accs = []
for epoch in range(50):

    start = time.time()
    epoch_train_loss, epoch_train_acc, gpu_mem = train_one_epoch(net, train_loader)
    epoch_test_loss, epoch_test_acc = evaluate_network(net, test_loader)
    epoch_val_loss, epoch_val_acc = evaluate_network(net, val_loader)

    print('Epoch {}, time {:.4f}, train_loss: {:.4f}, test_loss: {:.4f}, val_loss: {:.4f} \n
```

## Main function

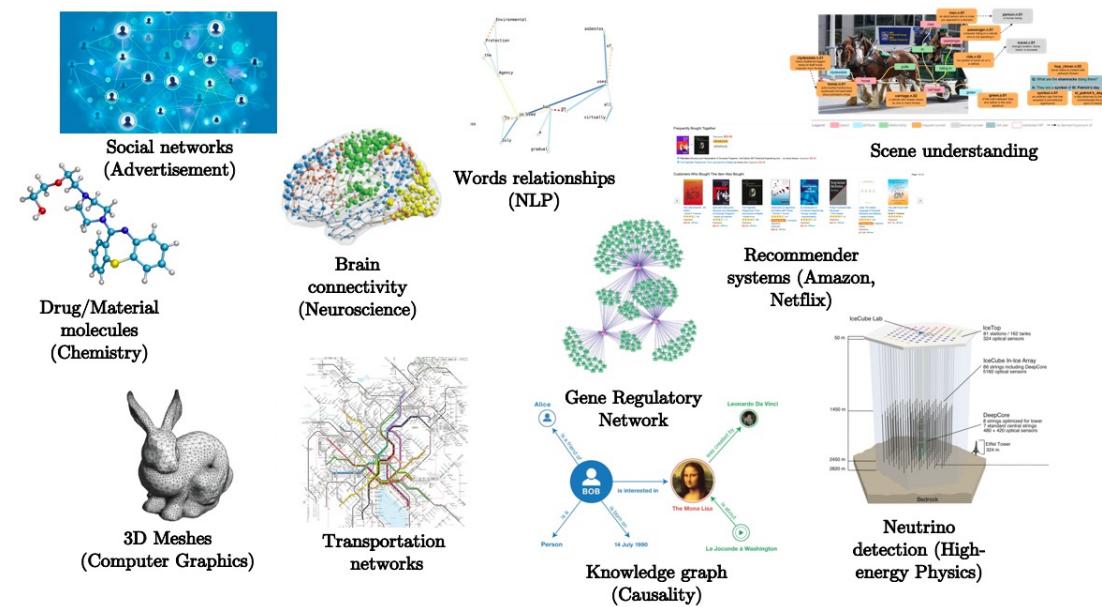
# Outline

- Part 1 : Traditional ConvNets
  - Architecture
  - Graph Domain
  - Convolution
- Part 2: Spectral Graph ConvNets
  - Spectral Convolution
  - Spectral GCNs
- Part 3 : Spatial Graph ConvNets
  - Template Matching
  - Isotropic GCNs
  - Anisotropic GCNs
  - DGL : Implemented layers and lab demo
- Conclusion

# Conclusion

- Contributions :
  - Generalization of ConvNets to data on graphs
  - Re-design convolution operator on graphs
  - Linear complexity for sparse graphs
  - GPU implementation (not yet optimized for sparse linear algebra multiplications)
  - Universal learning capacity
  - Multiple and dynamic graphs

- Applications :



“Graphs are the most important discrete models in the world!” -  
G. Strang (MIT)





Questions?